

Statistical Programming Language: South African Heart Disease Study

Due on March 30th, 2018

Alla Petukhina

Lukas Bargel (572482), Adrian Rolf (572656) and
Felix Vala (518231)

Table of Contents

List of Figures	ii
List of Tables	iii
1 Introduction	1
2 Methods	2
2.1 Fisher's Linear Discriminant	2
2.2 Logistic Regression	3
2.3 Leave-One-Out Cross-Validation	4
2.4 Cook's Distance	5
3 Descriptive Statistics	5
3.1 Description of Data Set	5
3.2 Graphical Exploration	7
4 Inferential Statistics	17
4.1 Fisher's Linear Discriminant	17
4.2 Logistic Regression	23
4.3 Leave-One-Out Cross-Validation	26
4.4 Cook's Distance	28
4.5 Results	33
5 Conclusion	36
References	iv
A Tables	v
B Code	vi
B.1 Descriptive Statistics	vi
B.1.1 Quantlet 1 - Data Preparation	vi
B.1.2 Quantlet 2 - Descriptive Statistics	viii
B.2 Inferential Statistics	x
B.2.1 Quantlet 3 - Fisher's Linear Discriminant	x
B.2.2 Quantlet 4 - Logistic Regression	xi
B.2.3 Quantlet 5 - Leave-One-Out Cross-Validation	xiii
B.2.4 Quantlet 6 - Cook's Distance	xiv
C Declaration Of Authorship	xvii

List of Figures

1	Correlation matrix of the different risk factors of coronary heart diseases.	11
2	Boxplots for different risk factors over occurrence of coronary heart disease	13
3	Jitterplot of selected outliers and the five most influential risk factors of coronary heart disease	16
4	Densities of the data according to the two classes of chd after projecting it onto an one-dimensional subspace using the Fisher's linear discriminant	22
5	The Cook's Distance values for every observation. The red line marks the threshold above which the observations are denoted as influential. Those observations are labeled with their id	31
6	The distributions of the covariates over chd . Highlighted are the observations that have been detected as influential with Cook's Distance	32

List of Tables

1	The summary statistics of the complete data set regarding the risk factors of coronary heart diseases	9
2	Selected outliers and their observations of the five most influential risk factors for coronary heart disease	14
3	The summary statistics of the five most influential risk factors of coronary heart disease for chd = 1.	14
4	The summary statistics of the logistic regression of typea with the explanandum chd	25
5	The summary statistics of the logistic regression of typea and age with the explanandum chd	25
6	The summary statistics of the logistic regression of m2	25
7	The out-of-sample performance for three different models from leave-one-out cross-Validation.	28
8	Results of the Logistic Regression models fitted on different sets of observations. All models use the same set of explanatory variables as chosen for m2 in section 4.2.	34
9	The summary statistics of the logistic regression of m0	v
10	The summary statistics of the logistic regression of m1.1	v
11	The summary statistics of the logistic regression of m1.2	vi
12	The summary statistics of the logistic regression of m3	vi

1 Introduction

Dying from a heart disease is still the leading cause for death worldwide. 15,5% of deaths in the world in the year 2015 (see World Health Organization, 2017) can be explained by ischaemic heart disease. As this is even stronger the case for upper-middle and high-income countries (18,3% and 16,9%), measures to prevent illness of the heart should be a number one priority for health care providers in those countries. To be able to design preventive measures to avoid heart disease from occurring, one should identify the major risk factors for this illness.

This paper revisits a classical article from a medical journal. In 1983, a group of researchers around J. E. Rossouw investigated the occurrence of heart disease in South Africa and collected data on life circumstances they considered to be risk-factors (see Rossouw et al., 1983). The aim of this paper is to reproduce the results of a classical empirical paper using current statistical methods. Thereby, capabilities of using of the statistical programming language *R* will be demonstrated. *R* is a language and environment for statistical programming and graphics. It is an evolution of the statistical language *S*, which was originally developed at AT&T's research and development by John Chambers and colleagues. *R* is very extensible as its open source nature provides the possibilities for developers to release extension packages, of which a few will be presented in this paper (cf. R-Foundation, 2018).

The paper is organized as follows. Section 2 gives a theoretical introduction to the statistical methods used in this paper. Section 3 explores the data set by using descriptive statistics. In section 4, the practical implementation of the methods presented in section 2 begins with the Fisher's linear discriminant in 4.1. In section 4.2, the Logistic Regression will be performed with different models and these will be evaluated based on their out-of-sample-performance in section 4.3 by implementing the leave-one-out cross-validation. As the last part of the analysis, the observations are examined for single outliers which heavily influence the model by Cook's distance in section 4.4. Furthermore section 4.5 summarizes the results of the previous chapters and section 5 describes the practical conclusions of the statistical results.

2 Methods

In the case at hand – the prediction if heart disease occurs or not given some crucial indicators – we are dealing with a binary classification problem. There are several methods that can be used for such problems.

2.1 Fisher’s Linear Discriminant

One convenient way is to use a linear discriminant or in particular, as we will do here, a Fisher’s linear discriminant.

Let $\mathbf{X} \in \mathbb{R}^{n \times p}$ denote a data matrix and $c \in \{0, 1\}$ a class, known a priori, that divides the data into two subsets \mathbf{X}_0 and \mathbf{X}_1 . The goal is now to find a separation rule that allows to allocate each observation – or new observations – into one of these subsets. The *best* separation rule is the one that minimizes the misclassification rate. Fisher’s approach to this minimization problem is, following the notation of Duda et al. (2001), to find a vector \mathbf{w} that projects each observation \mathbf{x}_i , $i = 1, \dots, n$ onto a one-dimensional subspace $y_i = \mathbf{w}^T \mathbf{x}_i$, that maximizes the ratio

$$J(\mathbf{w}) = \frac{|\tilde{m}_0 - \tilde{m}_1|^2}{\tilde{s}_0^2 - \tilde{s}_1^2}, \quad (1)$$

where $\tilde{m}_i = \frac{1}{n_i} \sum_{\mathbf{x} \in \mathbf{X}_i} \mathbf{w}^T \mathbf{x}$ denotes the sample mean and $\tilde{s}_i^2 = \sum_{\mathbf{x} \in \mathbf{X}_i} (\mathbf{w}^T \mathbf{x} - \tilde{m}_i)^2$ the scatter matrix of the projection of the subset \mathbf{X}_i , $i = 0, 1$. That is, finding a rule that keeps the *within-group-variance* as small as possible while keeping the *between-group-variance* as large as possible at the same time. The rule \mathbf{w} that maximizes equation 1 can be easily obtained by

$$\mathbf{w} = \mathbf{S}_w^{-1}(\mathbf{m}_0 - \mathbf{m}_1), \quad (2)$$

where

$$\mathbf{S}_w = \mathbf{S}_0 + \mathbf{S}_1, \quad (3)$$

with

$$\mathbf{S}_i = \sum_{\mathbf{x} \in \mathbf{X}_i} (\mathbf{x} - \mathbf{m}_i)(\mathbf{x} - \mathbf{m}_i)^T \quad (4)$$

is the sum of the scatter matrices of the non-projected data. \mathbf{w} is called the *Fisher's linear discriminant*.

The decision rule whether to allocate an observation to \mathbf{X}_0 or \mathbf{X}_1 now becomes a simple constant d that serves as a threshold. If we center the data first, this threshold is just 0. Thus, for each observation we finally get the discrimination rule:¹

$$\mathbf{w}^T \mathbf{x}_i \geq 0. \quad (5)$$

2.2 Logistic Regression

Another approach to binary classification problems, that comes with few assumptions and a representation as a model that can be easily used for interpretations is *Logistic Regression*. Let again $\mathbf{X} \in \mathbb{R}^{n \times p}$ denote a data matrix. The general idea of linear regression is to represent a *response variable* \mathbf{y} in terms of a linear function of a set of *explanatory variables* $\mathbf{x}_1 \dots \mathbf{x}_q$, $q \leq p$:

$$y_i = \beta_0 + \beta_1 x_{1i} + \dots \beta_q x_{qi} + \epsilon_i, \quad (6)$$

where $\epsilon_i \stackrel{\text{iid}}{\sim} \mathcal{N}(0, \sigma_\epsilon^2)$ is a stochastic error term. Taking the conditional mean of \mathbf{y} , the (zero mean) error term just drops out and we get, stated more conveniently in matrix notation:

$$\mathbb{E}[\mathbf{y} | \mathbf{X}] = \mathbf{X}\boldsymbol{\beta}. \quad (7)$$

An estimate for $\boldsymbol{\beta}$ can be obtained using Maximum Likelihood estimation and can be expressed analytically as $\hat{\boldsymbol{\beta}} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$. However, here we are assuming that $y_i \sim \mathcal{N}(\mathbf{x}_i^T \boldsymbol{\beta}, \sigma_\epsilon^2)$. For our binary classification problem, i.e. where $y_i \in \{0, 1\}$, this is not a very reasonable assumption. A more general approach to regression models that allow for several other distributions for y_i are *Generalized Linear Models*.

The idea of generalized linear models is to use a link function ϕ to express the model

¹Note that this problem can get quite more complex if non-equal prior probabilities and different distributions of the variables are involved. Thus, implicitly equal priors and a multivariate normal distribution are assumed here, making the calculation of \mathbf{S}_w and spotting of the threshold computational simple.

as

$$\phi(\mathbb{E}[\mathbf{y} | \mathbf{X}]) = \mathbf{X}\boldsymbol{\beta}. \quad (8)$$

It can be shown that there are several distributions – from the *exponential family* – that can be represented by a link function ϕ and that in such cases Maximum Likelihood estimation can be used to obtain estimates of the parameters (see e.g. Spokoiny and Dickhaus, 2015). If we choose $\phi(x) = \ln(\frac{x}{1-x})$ to be the *logit function* we obtain on the other hand

$$\mathbb{E}[\mathbf{y} | \mathbf{X}] = \phi^{-1}(\mathbf{X}\boldsymbol{\beta}) = \frac{e^{\mathbf{X}\boldsymbol{\beta}}}{1 + e^{\mathbf{X}\boldsymbol{\beta}}}, \quad (9)$$

where ϕ^{-1} is called the *logistic function*. This model is suitable to represent Bernoulli distributed response variables and is called a *Logistic Regression Model*. Estimates for $\boldsymbol{\beta}$ need to be obtained numerically. Fast algorithms for this purpose are for example the *Reweighted Least Squares Algorithm* (see e.g. Wasserman, 2013), which is the underlying algorithm in *R* for generalized linear models.

2.3 Leave-One-Out Cross-Validation

After fitting a model the question arises whether this model would still perform well on another dataset, sometimes referred to as the *out-of-sample performance*.² That is, to prevent over-fitting of the model to some particular dataset. But often, there is no other data at hand. One idea is then to split the data into a training and a test sample, the first to fit the model and the second to check its performance on the – from the view of the training sample – out-of-sample data. This can be repeated for different subsets and some measure over all the predictions can then be calculated, e.g. the *Mean Squared Error* or the misclassification rate. These measures can in turn be used as a benchmark for different models. This method is called *Cross-Validation*. We can choose the training sample to contain $n - 1$ of the n observation and the testing sample the remaining one. Repeating this procedure for every single observation is called *Leave-one-out Cross-Validation*.

²In the following, for simplicity, when we use this term, we refer to it as a measure of the out-of-sample performance obtained with leave-one-out cross-validation.

2.4 Cook's Distance

The *Cook's Distance* is a method to detect influential observations in regression models. It was proposed by Cook (1977) and is defined as

$$D_i = \frac{\sum_{j=1}^n \left(\hat{y}_j - \hat{y}_j^{(i)} \right)^2}{p \widehat{\text{MSE}}}, \quad (10)$$

where \hat{y}_j is the predicted value from our model for observation j , $\hat{y}_j^{(i)}$ the predicted value from our model for observation j excluding observation i and $\widehat{\text{MSE}} = \frac{1}{n-p} \sum_{i=1}^n (y_i - \hat{y})^2$ is an estimate of the mean squared error. As can be seen from equation (10), a high value of D_i compared to the values of the other variables indicates an influential observation i , since we would expect the deviations in the nominator to be approximately the same for all i if it doesn't matter whether we exclude it or not.

3 Descriptive Statistics

The following section deals with the discription of the data set (section 3.1) and the graphical exploration (section 3.2).

3.1 Description of Data Set

A postal campaign recruited 3357 white males and 3831 white females between the ages of 15 and 64 years as the population sample of the study. They represent 68% of the 1980 census population aged 15 – 64 years. Respondents completed a risk factor questionnaire and a London School of Hygiene questionnaire for chest pain by interview in a study centre set up in each community. The interviewers were dietitians who had been intensively trained in the standardized administration of the questionnaire. The questionnaire covered socio-economic items, smoking habits, a family history of heart disease, personal medical history and activity patterns. Topics like risk factor knowledge, attitudes and actions were covered as well. The respondents also completed a self-administered Bortner Short Rating Scale for

coronary-prone (**typea**) behaviour. (see Rossouw et al., 1983). This set-up lead to the following variables:

- To find out about the subject's heart condition, a resting 12-lead electrocardiography was recorded in the recumbent position on a Hewlett-Packard Model No. 1516 Tape Terminal for later playback and coding according to the Minnesota criteria. The coding was done in two stages and at each stage two independent blind observer were used, with an arbitrator for disagreements. This procedure lead to the dummy variable **chd**, which describes if the subject suffers from a heart disease or not.
- The variable **sbp** measures blood pressure. It was measured after the subject had been seated for 5 minutes. Following the American Heart Association guidelines for measuring blood pressure, a standard 12,5 x 23-cm cuff connected to a mercury manometer was used. Readings were taken three times and the lowest reading was recorded. Three observers were used to get safe results. Subsequent analysis of the data failed to show evidence of interobserver variation or of end-digit preference.
- **tobacco** observes tobacco usage in kilogramm. It measures the total tobacco consumed in the subject's lifetime and is calculated as the average per day multiplied by the period of use.
- The variable **typea** measures psychosocial stress on the self-administered Bortner Short Rating Scale.
- **famhist** is a dummy variable which describes if a family member of the subject suffered from a heart disease.
- **ldl** measures cholesterol levels. **ldl** is the amount of low-density lipoprotein. It is measured in mmol/l by taking and chemically analyzing³ blood samples of all participants.

³Boehringer CHOD-P AP enzymatic method and dextran sulphate-magnesium chloride precipitation (Rossouw et al., 1983)

- **obesity** measures the ration of weight to body-height of a person. This is commonly referred to as the so called body mass index of a person.
- **adiposity** measures the percentage of body fat of a person. While the most common measure for obesity would also be the body mass index, there are cases, where a high weight to height ratio is not caused by unhealthy body fat. Therefore, this measured variable gives us concrete information about the effect of excess body fat as a risk factor for heart disease.
- **age** measures the age of a person in years when the coronary heart disease occurs.
- The variable **alcohol** gives the information of individual alcohol consumption per year in liter.

All in all, the data set contains one response variable and nine potential risk factors for its occurrence. The following analyzed data set is a retrospective sample of 462 males, which was collected after some participants received blood pressure reduction treatment and other programs to reduce their risk factors.

3.2 Graphical Exploration

For a continuous workflow, we start by loading the necessary extension packages. The package `dplyr` will be used for data manipulation. `ggplot2` and `gridExtra` are packages with focus on data visualization. The package `corrplot` provides us with the function `corrplot()` for a correlation matrix. As the last package, `ElementStatLearn` includes the data set we us, namely `SAheart`. To create a robust code, it makes sense to be prepared for different scenarios. Additionally, warning messages in *R* are often quite difficult to understand for non-professionals. Therefore, we start by checking if all necessary packages are installed. `installed.packages()` gives a matrix which contains all packages. `"!"` defines as a logical NOT. So, if a package is missing, we formulated an easy understandable message for the user additionally to installing the package itself. With the option `character.only = TRUE` the `library()` command

interprets the passed argument as a string variable instead of searching for the word `pkg` in the library.

```

1  ## Check if required packages are installed, if not install and load
2  loadPKG <- function(pkg){
3    if (!(pkg %in% installed.packages())){
4      message(paste0("The required package '",
5                     pkg, "' is currently not installed. ",
6                     "It will now be installed to the default library path."))
7      install.packages(pkg)
8    }
9    library(pkg, character.only = TRUE)
10 }
11
12 loadPKG("dplyr")
13 loadPKG("ggplot2")
14 loadPKG("gridExtra")
15 loadPKG("ElemStatLearn")
16 loadPKG("corrplot")
17 loadPKG("xtable")
18
19 data(SAheart)

```

Code 1: Snippet of [Quantlet SAHeart_Q1_Data_Preparation](#)

Furthermore, we take care of the scenario when the PDF-path (save in folder desktop) does not exist. To check, we use `if()` in combination with `dir.exists()`, which gives back the information if the specified PDF-path exists. If that is not the case, the user is again informed with an easy understandable message and will be informed that the PDFs of the graphics will be saved in the current working directory. The name of the folder will appear in the message, as we use the function `getwd()` to show the name of the folder. `paste0` concatenates the strings without spaces in between.

```

22 ## Specify Path to save graphic outputs
23 PDFpath <- "~/Desktop"
24
25
26 ## check if PDFpath exists, if not choose current working directory
27 if (!dir.exists(PDFpath)){
28   PDFpath <- getwd()
29   message(paste0("The \'PDFpath\' that was specified does not exist. ",
30                  "Instead, the graphic outputs will be saved to the current ",
31                  "working directory. The current working directory is: \'",
32                  getwd(), "\'"))
33 }

```

Code 2: Snippet of [Quantlet SAHeart_Q1_Data_Preparation](#)

To get a first overview, we use the function `glm()`, which does the same as `str()` but shows as much data as possible. This is like a transposed version of `print`: columns run down the page, and data runs across. This makes it possible to see every column in a data frame (Wickham et al., 2017). The function `str()` is included in the code for comparison and the function `head()` for seeing the data in a table format. We also get information about the variable types. For comparison we included here both commands as well as the function `head()` that gives the data in an ordinary table format, where the number of the first rows to show may be specified as argument.

Additionally, we get a table of summary statistics by using the function `summary()`.

```
68 ## Some overview stuff
69
70 glimpse(SAheart)
71 str(SAheart)
72
73 head(SAheart, 5)
74
75 summary(SAheart)
```

Code 3: Snippet of [Quantlet SAHeart_Q1_Data_Preparation](#)

`summary()` is a useful generic function, which shows the results of various model fitting functions depending on the class of the object. As `SAheart` is a `data.frame`, we get a set of typical location parameters per variable (see table 1).

	Minimum	1 st Quantile	Median	Mean	3 rd Quantile	Maximum
sbp	101.0	124.0	134.0	138.3	148.00	218.0
tobacco	0.000	0.053	2.00	3.640	5.50	31.20
ldl	0.980	3.280	4.340	4.74	5.790	15.33
adiposity	6.740	19.77	26.11	25.41	31.23	42.49
typea	13.00	47.00	53.00	53.1	60.0	78.00
obesity	14.70	22.98	25.80	26.04	28.50	46.58
alcohol	0.000	0.510	7.510	17.04	23.89	147.2
age	15.00	31.00	45.00	42.82	55.00	64.00
famhist	Absent: 270 Present: 192					

Table 1: The summary statistics of the complete data set regarding the risk factors of coronary heart diseases (see related [Quantlet SAHeart_Q1_Data_Preparation](#)).

For example, we can immediately see that the age of the participants ranges from 15 to 64 years and there are less participants that have a family history of heart disease than have not. As *R* handles `famhist` as a factor variable, we have to rewrite it as an ordinary numerical variable, in order to calculate the correlation matrix using the `corr()` function, which is not able to handle factor variables. Factor variables are the way categorical variables are handled in *R* in various contexts and there are a lot of commands that can handle this variable type very utilitarian – including functions for regression models for example. On the other hand, there are also quite a lot of functions that cannot. Hence, it is advisable to convert to factor variables on the fly only when needed using the `as.factor()` function. Therefore, we decode `as.factor` as a dummy variable with 0 for absent and 1 for present. Technically, we refer to the underlying numerical values with the function `as.numeric()`. With the function

`head()` for the first ten values we can observe that present is defined as 2 and absent as 1. We convert the factor variable to a numerical variable with `as.numeric()` and redefine absent as 0 and present as 1.

```
78 ## Create id variable to easier handle single observations (using 'dplyr')
79 SAheart <- SAheart %>%
80   mutate(id = row_number())
81
82
83 ## famhist
84 head(SAheart$famhist, 10)
85 head(as.numeric(SAheart$famhist), 10)
86 SAheart$famhist <- as.numeric(SAheart$famhist)
87 SAheart$famhist[SAheart$famhist == 1] <- 0
88 SAheart$famhist[SAheart$famhist == 2] <- 1
```

Code 4: Snippet of [Quantlet SAHeart_Q1_Data_Preparation](#)

We start our exploratory analysis by creating a correlation matrix, which is a useful tool to get a systematic overview over correlations of the variables at a glance. From there, we get first intentions, which variables we should include in our model. We program the design of the matrix in such a way that intensity of color and shape of the graph show the strength and direction of the linear correlation between each of the variables. Technically, we therefore create a correlation matrix with `cor()` and use the `select()` function to put our response variable `chd` at the top. `ColorRampPalette()` gives us a vector of chosen colours, we use in the function `corrplot.mixed()`. This function gives us various options for the design of the matrix. We use upper and lower to give us shapes and colors in the upper diagonal field and numbers in the lower diagonal field.

```
59 ## Correlation matrix (using package corrplot)
60 pdf(file = paste0(PDFpath, "/Corrplot.pdf"), width = 7, height = 8)
61 corr_matrix <- cor(select(SAheart, chd, everything(), -id))
62 colset <- colorRampPalette(c("#BB4444", "#EE9988", "#FFFFFF", "#77AADD", "#4477AA"))
63 corrplot.mixed(corr_matrix,
64               upper = "ellipse",
65               lower = "number",
66               number.cex = .7,
67               upper.col = colset(10),
68               lower.col = "black",
69               tl.col = "black",
70               tl.pos = "lt")
71 dev.off()
```

Code 5: Snippet of [Quantlet SAHeart_Q2_Descriptive_Statistics](#)

The most striking finding (see figure 1) is that the further analysis should focus on the variables `sbp`, `tobacco`, `age`, `famhist`, `ldl` and `adiposity` as only they show noteworthy correlations with our explanandum `chd`. Furthermore, we observe a significant correlation between some of the explanatory variables. `adiposity` and `obesity` may be a result of a similar nutrition, so the resulting variables are strongly

correlated. As the last obvious point, **obesity** and **adiposity** are correlated with **age**.




Figure 1: Correlation matrix of the different risk factors of coronary heart disease (see related [Quantlet SAHeart.Q2.DescriptiveStatistics](#)).

To further examine these first results, we program a number of boxplots. As we want to look at the distribution of the variables over 0 and 1 in **chd** we concentrate only on the five variables with high correlations as we don't expect distinguishable differences between the boxplots for 0 and 1 for the variables with low correlations. The Boxplots are created with the self-written function **multiBP**, that takes a list of variables that shall be plotted over **chd** as input. Optionally, a path can be specified to save the graphic in pdf format to the local machine. We use the function **lapply** to perform the generation of the plots on each of the variables contained in the input list. **lapply** in turn returns a list of specific **ggplot** graphic objects. **ggplot** is a powerful toolkit to handle and generate graphics within in *R* from data. Besides of specifying the underlying dataset and the variables to be used – where we convert **chd** on the fly to a factor variable – **ggplot()** takes various options and so called *layers* that can be connected via a "+" symbol. The different layers make it easily possible to merge

different illustration of the same data into one graphic. Here we only use one layer for the boxplots. Also, the appearance of the graphic can be edited in great detail or just templates - like `theme_bw()` – can be used, as we do here. Since the number of boxplots that will be saved to `plot.list` can vary, we account for the arrangement of the single boxplots in one graphic. If there are more than three boxplots, the number of columns is restricted to three, so that the `grid.arrange()` function will successively push the fourth and seventh into the next rows. Otherwise all boxplots would be displayed in one row and therefore heavily shrunk in their width. `grid.arrange()` as well as `ggsave` are functions specially designed to handle `ggplot` graphic objects.

```

74 ## Function to show Boxplots of covariates over chd
75 multiBP <- function(varlist, PDFpath = NULL){
76   plot.list <- lapply(varlist, function(ivar){
77     ggplot(data = SAheart, mapping = aes(x = as.factor(chd), y = get(ivar))) +
78       geom_boxplot(stat = "boxplot", position = "dodge") +
79       theme_bw() +
80       xlab("chd") +
81       ylab(ivar)
82   })
83
84   if(length(plot.list) <= 3){
85     ncol <- length(plot.list)
86   } else {
87     ncol <- 3
88   }
89
90   final.plot <- grid.arrange(grobs = plot.list, ncol=ncol)
91
92   if(!is.null(PDFpath)){
93     ggsave(filename = "/BPplot.pdf", path = PDFpath, device = "pdf", plot = final.plot)
94   }
95 }
96
97 varlist <- c("sbp", "tobacco", "age", "ldl", "adiposity")
98 multiBP(varlist, PDFpath)

```

Code 6: Snippet of [Quantlet SAHeart.Q2.Descriptive_Statistics](#)

All boxplots solidify our assumptions (see figure 2), that the five variables might be important contributors for our model. The boundaries of all quantiles lie higher for observations where the participants of the study suffered from heart disease. Especially `age` and `tobacco` show a higher median at `chd = 1` than where the 75% quantile of the values with `chd = 0` lies. Boxplots also give first signs for outliers. Observations, which are further away from the box than 1,5 times the interquartile range, are shown as dots. We can observe an especially noteworthy outlier in the boxplot graphic belonging to the variable `age` as it is the only outlier to the bottom. As we take a closer look at the observation, we can see that the person has quite – medically spoken – unobtrusive values for all risk factors and suffers from heart disease nevertheless. We keep that in mind for our model selection later. Furthermore, we observe relatively strong outliers on `sbp` and `ldl`. Despite the thorough

procedure of measurements, it is possible, that these are measurement errors since the occurrence of such high values in the human body is unlikely.

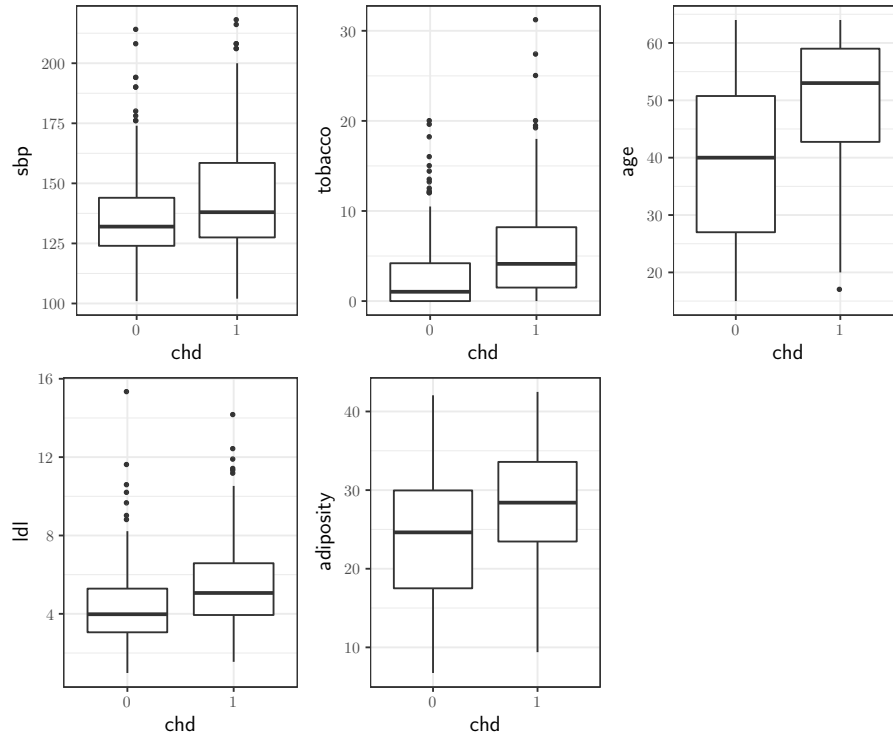


Figure 2: Boxplots for different risk factors over occurrence of coronary heart disease (see related [Quantlet SAHeart_Q2_Descriptive_Statistics](#)).

As the next part of our analysis, we want to observe the behavior of selected outliers. From the boxplots we know, that there is a very low value of **age** at **chd** = 1. So the scenario we look at in the following are very young participants (defined as **age** \leq 20), who suffer from heart disease. To get the IDs of these observations, we let *R* show us the values of the variable ID, but only those with the restrictions that the variable **age** is smaller or equal 20 and **chd** is equal to 1. With the "&" sign we define that both conditions have to be met.

```

102 ## Suspicious outlier
103 SAheart$id[SAheart$age <= 20 & SAheart$chd == 1]
104
105 subset(SAheart, select = varlist, id == 261)
106 subset(SAheart, select = varlist, id == 21)
107 summary(subset(SAheart, select = varlist, chd == 1))

```

Code 7: Snippet of [Quantlet SAHeart_Q2_Descriptive_Statistics](#)

As a result, we get to know that the conditions are met by the observations with IDs 21 and 261. To further examine these two observations, we use the function

`subset()` to show the values (see table 2) of these specific variables.

	261	21
sbp	118	106
tobacco	0.00	1.61
age	17	20
ldl	2.39	1.74
adiposity	12.13	12.32

Table 2: Selected outliers and their observations of the five most influential risk factors for coronary heart disease (see related [Quantlet SAHeart_Q2_Descriptive_Statistics](#)).

For comparison, we create some summary statistics (see table 3) with the function `summary()` as before, but this time only for the five most influential risk factors and for observations where `chd = 1`. A first result is that our two selected observations have unusually low values for nearly all variables.

	Minimum	1 st Quantile	Median	Mean	3 rd Quantile	Maximum
sbp	102.0	127.5	138.0	143.7	158.5	218.00
tobacco	0.000	1.500	4.130	5.525	8.200	31.200
age	17.00	42.75	53.00	50.29	59.00	64.000
ldl	1.550	3.940	5.065	5.488	6.582	14.160
adiposity	9.390	23.46	28.41	28.12	33.59	42.490

Table 3: The summary statistics of the five most influential risk factors of coronary heart disease for `chd = 1` (see related [Quantlet SAHeart_Q2_Descriptive_Statistics](#)).

To visualize these results, we create one jitterplot per variable, where the two selected outliers are highlighted in color. That is reasonable since we can see a first glance of the behavior of the highlighted observations in comparison to all others on the selected variables. For this purpose, serves the function `multiJP`, which takes `varlist` as an obligatory argument. Optionally, there may be specified a list of observations via IDs to be highlighted, and if they shall be labeled with their IDs. And again, it is possible to specify a path to save the graphic in pdf format to the local computer. First, the function checks if the Boolean parameter `label` was set to `TRUE`. If so, the `varlist` is expanded by an additional entry, indicating that a legend should be set. Like above, this `varlist` is passed to an `lapply()` function that will return a

list containing the single jitterplots. To obtain a legend that gives the IDs of the highlighted observations we use a workaround here – that is why the corresponding variable is called `hackdata` – since a proper legend can’t easily be created with the standard function in this special case here.

```

110 multiJP <- function(varlist, obslist = 0, PDFpath = NULL, label = FALSE){
111
112
113   if(label){
114     varlist <- c(varlist, "legend")
115   }
116
117   color.list <- c("firebrick", "dodgerblue", "forestgreen",
118                 "orangered", "violet", "yellow3")
119
120   plot.list <- lapply(varlist, function(ivar){
121     if(ivar == "legend"){
122       hackdata <- data.frame(id = obslist,
123                             position = seq(1,length(obslist)),
124                             xvalue = 0.11)
125       legend <- ggplot(data = hackdata, mapping = aes(x = xvalue, y = position)) +
126         geom_point(colour = color.list[1:length(obslist)], shape = 17, size = 3) +
127         theme_classic() +
128         theme(axis.text = element_blank(), axis.ticks = element_blank(),
129               axis.line = element_blank()) +
130         labs(x = "", y = "") +
131         scale_y_continuous(expand = c(0, 0), limits = c(0, max(hackdata$position)+1)) +
132         scale_x_continuous(limits = c(0.1, 1)) +
133         geom_text(data = hackdata, aes(label = id),
134                  colour = color.list[1:length(obslist)], size = 3, hjust = -1)
135     } else {
136       ggplot(data = SAheart[!(SAheart$Id %in% obslist),],
137             mapping = aes(x = as.factor(chd), y = get(ivar))) +
138         geom_jitter(shape = 1) +
139         theme_bw() +
140         xlab("chd") +
141         ylab(ivar) +
142         geom_jitter(data = SAheart[SAheart$Id %in% obslist,],
143                    colour = color.list[1:length(obslist)], shape = 17, size = 3) # +
144     }
145   })
146
147   if(length(plot.list) <= 3){
148     ncol <- length(plot.list)
149   } else {
150     ncol <- 3
151   }
152
153   final.plot <- grid.arrange(grobs = plot.list, ncol=ncol)
154
155   if(!is.null(PDFpath)){
156     ggsave(filename = "/Jitterplot-outliers.pdf", path = PDFpath, device = "pdf", plot = final.plot)}
157

```

Code 8: Snippet of [Quantlet SAHeart.Q2.Descriptive.Statistics](#)

The idea is to add an additional graphic that is build such that it serves as a legend. This is obtained by removing all elements like axes or gridlines and only displaying the highlighted observations in the corresponding color and their ids attached to them as labels. All necessary information is saved to `hackdata` as a `data.frame`. That is, the IDs as well as the positions of the symbols – depending on the number of observation – that shall appear in the legend. While scheduling the variables within the `lapply()`-procedure we check at each step if the last entry *legend* is reached. If not, a jitterplot for the current variable is created similar as the boxplots are created within the `multitBP()` function described above. If instead the entry *legend* is reached

– given of course it was specified – the legend graphic will be build. This way `lapply()` returns a full list containing the jitterplot as well as the legend graphic. This is a crucial point here. Once the list was returned it is not possible to add additional graphics here since this list is already stored as a specific `ggplot` object ready to be processed further by specific functions, namely `grid.arrange` and `ggsave()` in our case. To use the function for our concrete problem, we define `varlist` as a vector of all variables, which we determined as relevant in the correlation matrix and `obslist` as a vector of the two selected outliers with IDs 21 and 261.

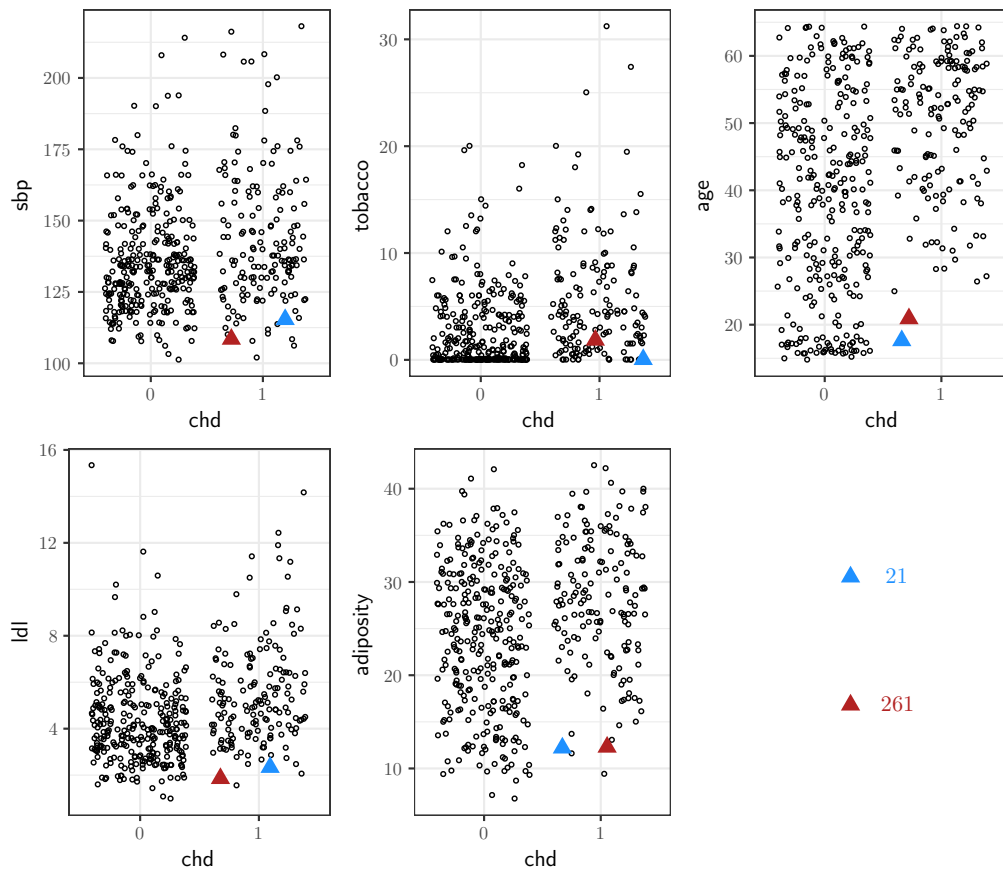


Figure 3: Jitterplot of selected outliers and the five most influential risk factors of coronary heart disease (see related [Quantlet SAHeart_Q2_Descriptive_Statistics](#)).

In the resulting graphic (see figure 3) we can observe that the two selected outliers have very low values on all variables. We keep that in mind for our model selection later. Very young persons who nevertheless suffering from heart disease are the most obvious possibility for outliers, but not the only one. To generally check for further conspicuous cases, we search for observations, which have very low values

throughout all variables and anyhow suffer from heart disease (We check for low values, since they stand for a healthier lifestyle or body condition in all variables, so one wouldn't expect the occurrence of heart disease). Hence, we filter for cases where `chd = 1` and all other variables lie in the 25% quantile. We let `R` give us the IDs of these observations and can see, that this scenario applies to none of the observations. Hence, in this case no further graphical examination is needed. However, we will make use of this function again in section 4.4.

4 Inferential Statistics

This section begins with the practical `R`-implementation of the methods presented in section 2 by coding the Fisher's linear discriminant in 4.1. In section 4.2, the Logistic Regression will be performed with different models and these will be evaluated based on their out-of-sample-performance in section 4.3 by implementing the leave-one-out cross-validation. As the last part of the analysis, the observations are examined for single outliers which heavily influence the model by Cook's distance in section 4.4. Furthermore section 4.5 summarizes the results.

4.1 Fisher's Linear Discriminant

In this section we describe the implementation of the function for the *Fisher's linear discriminant*. First, we calculate the Fisher's linear discriminant \mathbf{w} as described in section 2.1. Then, we apply the decision rule from equation (5) to the data and calculate the prediction accuracy as $1 - \text{missclassification rate}$, i.e. the proportion of correct predictions. Further, we will plot the densities of the projected data according to the two classes to visualize how good this decision rule is able to difference between these two classes.

The function `fisher` takes three arguments. First, `X`, which can be a matrix or a `data.frame` object of dimension $(n \times q)$, that contains the covariates, i.e. the set of variables that shall be used to classify into one or the other class. Second, `c`, a binary vector of dimension $n \times 1$, containing the respective class for every observation. Note, that both arguments are obligatory since no default is set. And, third, an optional

parameter `PDFpath` that can, again, be used to directly save the graphical output – in this case the densities of the projected data given the classes – to the local computer.

In line 54 we use the `as.matrix()` function to convert the input into a matrix in case it is given as a `data.frame` object, since the operations in the following calculations require matrix objects. We then split the matrix into two matrices according to their classes using simple index notation. Leaving the right side of the comma blank means that no condition for the single columns is stated and therefore refers to all columns. The command `nrow()` returns the number of rows. In lines 256 and 257 we calculate the mean for each column of the two matrices and convert them to a vector to get a plain numeric vector without additional information like the variable names saved to this object.

```

51 ## Function for Fisher Linear Discriminant
52 fisher <- function(X, c, PDFpath = NULL){
53
54   X <- as.matrix(X)
55
56   X0 <- X[c == 0,]
57   X1 <- X[c == 1,]
58
59   n0 <- nrow(X0)
60   n1 <- nrow(X1)
61
62   m0 <- as.vector(colMeans(X0))
63   m1 <- as.vector(colMeans(X1))
64
65   M0 <- as.vector(rep(1, n0)) %*% t(m0)
66   M1 <- as.vector(rep(1, n1)) %*% t(m1)
67
68   X0_c <- X0 - M0
69   X1_c <- X1 - M1
70
71   S0 <- t(X0_c) %*% X0_c
72   S1 <- t(X1_c) %*% X1_c
73
74   S_w <- S0 + S1
75
76   S_w_inv <- solve(S_w)
77
78   w <- S_w_inv %*% (m0 - m1)

```

Code 9: Snippet of [Quantlet SAHeart_Q3_Fisher_Linear_Discriminant](#)

This prevents errors during matrix operations. Not converting it sometimes causes the error, that *R* doesn't recognize it as a proper numeric vector. However, the results are two $(q \times 1)$ vectors.

In the next step we want to center the matrices \mathbf{X}_i using the respective mean vectors \mathbf{m}_i , $i = 0, 1$. What is mathematically easily expressed as $\mathbf{X}_i - \mathbf{m}_i$ needs a little more work in *R*. *R* can't handle operations like $\mathbf{Z}_{(m \times n)} - \mathbf{z}_{(m \times 1)}$. The operands need to have the same dimension. For this purpose we want to create matrices that contain n times the q means for each variable in the columns. This can be achieved by

$\mathbf{1}_{(n_i \times 1)} \mathbf{m}_i^T$, where $\mathbf{1}_{(n_i \times 1)}$ is a vector of ones. $\mathbf{1}_{(n_i \times 1)}$ is created using the function `rep()` that returns a sequence of ones of length n_i in this case. This sequence is then converted to a vector so R can easily handle it during matrix operations.

The `%*%` denotes the matrix multiplication operation and the function `t()` transposes a matrix. With these build-in operations of R we can calculate the \mathbf{M}_i and finally use them to center the matrices \mathbf{X}_i using simple subtraction.

The same goal could have of course been achieved by using a loop running over each row of \mathbf{X}_i subtracting \mathbf{m}_i . But the larger the covariates matrix \mathbf{X} the more computational effort is needed to run over every single row. Matrix operations are very efficient implemented in the most programming languages, including R. Hence, whenever possible, programmed algorithms should use matrix operations instead of iterations to avoid time consuming calculations.

Using the centered, data we calculate the sum of the scatter matrices as described in equation (3) and (4). To get the inverse of \mathbf{S} we use the function `solve(a, b)`, where `a` and `b` are two arguments. The purpose of this function is to solve a system of equations that is given in terms of matrices and vectors $\mathbf{Ax} = \mathbf{b}$. If the argument `b` is not specified an identity matrix \mathbf{I} of the respective dimension according to \mathbf{A} is assumed as default. Thus, the equation becomes $\mathbf{Ax} = \mathbf{I}$ and the solution to this equation is \mathbf{A}^{-1} . Therefore, this function can easily be used to obtain the inverse of a matrix – given it is non-singular, of course. Finally, we have all components together to calculate the Fisher’s linear discriminant \mathbf{w} .

```

80 m <- as.vector(colMeans(X))
81 M <- as.vector(rep(1, nrow(X))) %*% t(m)
82 X_c <- X - M
83
84 z <- t(w) %*% t(X_c) < 0
85
86 accuracy <- sum(t(z) == c) / length(z)

```

Code 10: Snippet of [Quantlet SAHeart_Q3.Fisher.Linear.Discriminant](#)

Afterwards we center the whole covariates matrix X as above. The projection of the centered data is then used for the decision rule described in equation (5). To apply the decision rule on the whole covariates matrix at once instead for every single observation, \mathbf{X}_c^T is used – since we achieve the same goal by $\mathbf{w}^T \mathbf{X}_c^T$ as by calculating $\mathbf{w}^T \mathbf{x}_i$ n -times. If the value of the projected data is lower than zero the $(n \times 1)$ vector \mathbf{z} will contain TRUE for this row and FALSE otherwise. Thus, \mathbf{z} consists of

the predictions to which class an observation is allocated by the decision rule, where `TRUE` refers to class 1. That is in our case the prediction that an observation suffers from coronary heart disease.

Furthermore, the prediction accuracy is calculated. `z` is transposed here so that `R` can check for equality on each row of `z` and `c` in one step. This comparison is possible here since the 1 and 0 entries of `c` are automatically handled as `TRUE` and `FALSE`, respectively. Each comparison in turn yields `TRUE` or `FALSE` such that summing this up is just like counting the number of `TRUE`-entries. Dividing by the total number of rows, we finally get the prediction accuracy. Note, that we use `length()` here since it returns n in both cases, whether the input is a $(n \times 1)$ or an $(1 \times n)$ vector. Otherwise, depending on the dimension in which the vector is currently given within `R`, one has to check whether `nrow()` or `ncol()` is appropriate. Thus, in such cases `length` is the better choice to avoid mistakes that can easily be overseen at first glance.

```

88 Y0 <- X0 %*% w
89 Y1 <- X1 %*% w
90
91 Y <- rbind( cbind(Y0, rep(0, nrow(Y0))),
92            cbind(Y1, rep(1, nrow(Y1))) )
93 Y <- as.data.frame(Y)
94 colnames(Y) <- c("Y", "c")

```

Code 11: Snippet of [Quantlet SAHeart.Q3.Fisher.Linear.Discriminant](#)

In order to plot the densities of the projected data according to the classes, we first need to calculate the respective vectors containing these data. This is done in lines 88 to 89. For each of the two vectors an additional column is added containing the respective class using `cbind()`, a command to easily bind columns of vectors and/or matrices together. Again, we use the `rep()` command to create a vector of ones or zeros, respectively, of the right length on the fly. After that, these two matrices are binded again, this time horizontally using `rbind()`, yielding in the matrix `Y` containing the projection of every observation and the respective class. Furthermore, we need to convert `Y` into a `data.frame` object, since the following `ggplot()` command can only handle this type of objects. At last, we change the column names of `Y` in order to easily refer to its columns.


```

96 final.plot <- ggplot(Y, aes(Y, fill = as.factor(c))) + geom_density(alpha = 0.2) +
97   theme_classic() +
98   labs(x = "", y = "") +
99   theme(axis.text.x = element_blank(), axis.ticks.x = element_blank(),
100         axis.text.y = element_blank(), axis.ticks.y = element_blank(),
101         axis.line.y = element_blank()) +
102     scale_fill_manual(values = c("firebrick", "dodgerblue"),
103                      name = "chd", labels = c("no", "yes"))
104
105 plot(final.plot)
106
107 if (!is.null(PDFpath)){
108   ggsave(filename = "/Densitiesplot.pdf", path = PDFpath, device = "pdf", plot = final.plot)
109 }
110
111 return(accuracy)
112 }
113
114
115 X <- select(SAheart, -id, -chd)
116 c <- SAheart$chd
117
118 fisher(X=X, c=c, PDFpath)

```

Code 12: Snippet of [Quantlet SAHeart_Q3 Fisher Linear Discriminant](#)

Furthermore, we create the plot of the densities of the projected data according to the two classes. The basic parameters of `ggplot()` have already been described in section 3.2. Here, it can be seen that using `Y` as a data.frame that contains the projected data as well as the respective classes makes it easy to use and the code easy to read. With `geom_density()` we define the type of the graph to be a density in this case. That is, Kernel density estimates are calculated for the given data, where a gaussian kernel is used by default. The parameter `alpha = 0.2` makes the areas of the distributions transparent, so we can also see the overlapping areas.

The additional parameters specified within the `theme()` command change the look of the graphic, in particular removing the ticks of the axis as well as the labels and the line of the y axis. The specific values are not of interest here since we are only interested in how good the two densities are separated from each other. If they had no overlapping areas that would be the best result since then we could perfectly allocate any observation to the right class with a probability of zero for misclassification. But this is usually not the case.

With `scale_fill_manual` the color of the areas of the density is specified and at the same time the appearance and labels of the legend is configured. Finally, the graphic is plotted and, as already described in section 3.2, saved to the local machine if the `PDFpath` argument was specified. The prediction accuracy, i.e. a single value, is returned by the `fisher` function.




Figure 4: Densities of the data according to the two classes of **chd** after projecting it onto an one-dimensional subspace using the Fisher's linear discriminant (see related [Quantlet SAHeart_Q3.Fisher.Linear.Discriminant](#)).

After selecting all variables, except of **chd** and **id**, as the covariates matrix \mathbf{X} and **chd** as the class vector \mathbf{c} , we run the fisher function. A prediction accuracy of ca. 0.695 is returned. That is a relatively satisfying result for such a case. At least it is better than a random classification of 0.5, but has still a misclassification rate of over 30%. The reason for that can be seen in Figure 4. The two densities have quite a large overlapping area. This means that in this cases we decide for the class that has a higher probability but there is still non-zero or even only a marginally less high probability to belong to the other class. Thus, there are a lot of cases where this decision rule might allocate an observation to the wrong class.

The Fisher's linear discriminant method has some weaknesses. We assume implicitly that every variable from the covariates matrix is normal distributed and this is of course not necessarily the case and especially for the binary variable **famhist** this assumption is obviously wrong. Anyway, Figure 4 already gives us a clue that we might not find a decision rule that can allocate the observations with a misclassification rate lower than 20% or even 10% at all. However, we can still try to improve the classification. In the next section we will use logistic regression models to formulate a decision rule. Logistic regression needs few assumptions – e.g. no

assumptions about the distributions of the covariates – and is known to be a robust method for binary response variables (Wasserman, 2013).

4.2 Logistic Regression

In the following we will use logistic regression to examine possible models. In line with the methodical section 2 we are using the generalized linear model function `glm()` from the package `stats` loaded in *R* by default. This function takes a regression model as an obligatory argument. The syntax is *dependent var : independent var1 + independent var 2 + .* Also, the underlying dataset has to be specified if the variable path is not provided in detail – i.e. like *somedataset\$somevar* – or the respective dataset hasn't been attached to the *R* environment. Furthermore, we have to set the `family` parameter to indicate which specific case of a generalized model we are dealing with. In our case we have to set a `binomial` family and additionally the link to `logit`, see section 2.2 for more details (note that the Bernoulli distribution is a special case of the Binomial distribution, so this refers to the same distribution here). Again, we use the `summary()` command to obtain the regression outputs. As already mentioned in section 3.2, it can handle various models and provide suitable outputs.

```
121 ## M0: (naive) Full model
122 m0 <- glm(chd ~ ., family = binomial(link = "logit"),
123          data = select(SAheart, -id))
124 summary(m0)
125
126
127 ## M1: Without sbp, alcohol; and exclude either adiposity or obesity
128 m1.1 <- glm(chd ~ tobacco + age + famhist + ldl + typea + adiposity,
129            family = binomial(link = "logit"), data = SAheart)
130 summary(m1.1)
131
132 m1.2 <- glm(chd ~ tobacco + age + famhist + ldl + typea + obesity,
133            family = binomial(link = "logit"), data = SAheart)
134 summary(m1.2)
135
136
137
138 ## M2: Exclude sbp, alcohol, adiposity, obesity
139 m2 <- glm(chd ~ tobacco + age + famhist + ldl + typea,
140          family = binomial(link = "logit"), data = SAheart)
141 summary(m2)
142
143
144 ## M3: For sake of simplicity of models we try to exclude
145 ## typea, since effect seems to be still very low
146 m3 <- glm(chd ~ tobacco + age + famhist + ldl,
147          family = binomial(link = "logit"), data = SAheart)
148 summary(m3)
```

Code 13: Snippet of [Quantlet SAHeart_Q4_Logistic_Regression](#)

Despite the fact that we already received first hints on influential variables from the graphical exploration, we start with a naive full regression model `m0`, i.e. including

all variables. Thus, we possibly find some effects that get visible while controlling for other variables, since plain correlations do not account for that. A full model can easily be obtained using the single point shortcut, which serves as a placeholder for all variables of the specified dataset, except of the dependent variable, in this case. The output (see table 9) shows that **sbp** and **alcohol** have quite high p-values. Therefore, they are excluded in the next models.

We already know from theory as well as from figure 1 that **obesity** and **adiposity** are highly correlated. Hence, we perform two models **m1.1** and **m1.2** excluding **obesity** or **adiposity**, respectively. The outputs (for **m1.1** see table 10 and for **m1.2** see table 11) show that in both models the respective variable has a high p-values. Therefore, we exclude both variables and obtain the model **m2**, where all explanatory variables are significant at least at a 0.1% significance level. With regard to the relatively low correlation of **typea** with **chd** (see Figure 1), it was reasonable to exclude **typea** from **m3**. Nevertheless, the AIC of **m3** (see table 12) is worse than the AIC of **m2** and the estimated coefficient of age in **m3** is slightly weaker. Interestingly, **typea** has a significant effect, what we would not have expected from the graphical exploratory analysis. For sake of simplicity we exclude **typea** in model **m3**. This increases the AIC by more than 7 points compared to model **m2**.

```

152 ## typea: mediator effect via age
153
154 mean(SAheart$chd[SAheart$typea >= median(SAheart$typea)])
155 mean(SAheart$chd[SAheart$typea < median(SAheart$typea)])
156
157 multiBP("typea")
158
159 obslist <- (SAheart %>%
160   filter(chd == 1 & typea < 30) %>%
161   select(id))$id
162
163 varlist <- c("sbp", "tobacco", "age", "ldl", "typea", "famhist")
164 multiJP(varlist, obslist, label = TRUE)
165
166
167 summary(glm(chd ~ typea, family = binomial(link = "logit"), data = SAheart))
168 summary(glm(chd ~ typea + age, family = binomial(link = "logit"), data = SAheart))

```

Code 14: Snippet of [Quantlet SAHeart.Q4.Logistic-Regression](#)

Technically, we compare the mean and median of **typea** as well as generating a boxplot with the function **multiBP()** to get a rough threshold, which we use for the jitterplot **multiJP** to identify outliers. Moreover, we establish two controlling-models in the **glm()** environment. The first one consists of **typea** and the second one adds **age**. Comparing a logistic regression only containing **typea** as explanatory variables

with a model where we control for **age** shows, that the p-value drastically decrease in the second case (see table 4 and 5).

	Estimate	Std. Error	z value	Pr(> z)	
(Intercept)	-1.8447	0.5607	-3.29	0.0010	**
typea	0.0226	0.0103	2.20	0.0275	*
Significance codes	0 \Rightarrow ***	0.001 \Rightarrow **	0.01 \Rightarrow *		
AIC	595.12				

Table 4: The summary statistics of the logistic regression of **typea** with the explanandum **chd** (see related [Quantlet SAHeart.Q4.Logistic.Registration](#)).

	Estimate	Std. Error	z value	Pr(> z)	
(Intercept)	-5.9170	0.8414	-7.03	0.0000	***
typea	0.0396	0.0115	3.44	0.0006	***
age	0.0702	0.0091	7.75	0.0000	***
Significance codes	0 \Rightarrow ***	0.001 \Rightarrow **	0.01 \Rightarrow *		
AIC	519.04				

Table 5: The summary statistics of the logistic regression of **typea** and **age** with the explanandum **chd** (see related [Quantlet SAHeart.Q4.Logistic.Registration](#)).

Thus, we can conclude that **age** serves as a cofounder here regarding **typea** and **chd**. This is sometimes called an *incidental cancellation*. Therefore, we will keep **typea** and select **m2** as the best model here, since it has the lowest AIC value as well as the highest significant coefficients, which can be regarded at table 6.

	Estimate	Std. Error	z value	Pr(> z)	
(Intercept)	-6.4464	0.9209	-7.00	0.0000	***
tobacco	0.0804	0.0259	3.11	0.0019	**
age	0.0505	0.0102	4.94	0.0000	***
famhist	0.9082	0.2258	4.02	0.0001	***
ldl	0.1620	0.0550	2.95	0.0032	**
typea	0.0371	0.0122	3.05	0.0023	**
Significance codes	0 \Rightarrow ***	0.001 \Rightarrow **	0.01 \Rightarrow *		
AIC	487.69				

Table 6: The summary statistics of the logistic regression of **m2** (see related [Quantlet SAHeart.Q4.Logistic.Registration](#)).

4.3 Leave-One-Out Cross-Validation

After calculating different possible models, we now want to check their out-of-sample performance as described in section 2.3 as another measure for the quality of the models and hence, for model selection.

For this purpose we wrote the function `loo_cv`, which takes two arguments. The first, `model` is an obligatory argument that needs to be a `glm` or `lm` object. These kind of objects are returned by the eponymous functions as introduced in section 4.2. The second, `cutoff` is an optional argument that marks the threshold whether an observation shall be allocated to class 0 or 1, since the prediction for an observation from a logistic regression model can be interpreted as a probability. The default is 0.5, i.e. if the probability to suffer from coronary heart disease is higher than 0.5 this observation will be allocated to class 1.

The underlying dataset of the input model and the formula of the regression model are read from the `glm` object. These objects are handled like a `data.frame`. Thus, we can easily access different information simply using the `$`-notation. Generally, this makes these objects to be easily used by other functions, e.g. from different packages, that perform tests, forecasting, interval calculations, graphical representations, etc. Furthermore, the number of rows is calculated and an empty vector is created using the `vector` command. Here, the `mode`, that is the type of the content, and the `length` are specified. The reason for this empty vector is that in the next step, we will write to it row by row. If the vector is already specified, we can use simple index notation for that. Of course, this is only possible in cases where we know the number of rows – or more generally entries – that will be written to the vector in advance. Another possibility is to create a vector or a list that will be appended on the fly in each iteration step. However, the first solution is recommendable because if the resulting vector is used in further operations its length/dimension might be important for functionality, e.g. if matrix multiplication were used. This could cause errors that might be difficult to track. Knowing the length/dimension allows to write better assessable upcoming code.

Afterwards, we are running over each observation fitting the model that was given into the `loo_cv` function using a `for()` loop. Note that the expression `1:N` creates a

sequence $1, \dots, N$ on the fly from which we take every element to iterate over, saved to the variable `i` in each step, so we can refer to it. In general, we could specify an arbitrary list instead of `1:N` to iterate over each of its elements, e.g. the letters of the alphabet.

```

72 loo_cv <- function(model, cutoff = 0.5){
73
74   cur.data <- model$data
75   cur.formula <- model$formula
76   N <- nrow(cur.data)
77
78   prediction <- vector(mode = "numeric", length = N)
79
80   for(i in 1:N){
81     cur.model <- glm(cur.formula, family = binomial(link = "logit"),
82                     data = cur.data[-i,])
83     prediction[i] <- predict(cur.model, newdata = cur.data[i,], type = "response")
84   }
85
86   correct <- cur.data$chd == (prediction > cutoff)
87
88   accuracy <- sum(correct)/N
89
90   return(accuracy)
91 }
92
93
94 loo_cv(m0)
95 loo_cv(m2)
96 loo_cv(m3)

```

Code 15: Snippet of [Quantlet SAHeart_Q5-Leave-One-Out-Cross-Validation](#)

Again, we use the `glm()` function as described in section 4.2. The formula that was read from the input model is inserted here as well as the data. In each iteration step the observation i is removed from the dataset and the model fitted without it. Again, this is easily achieved using index notation. With `-i` a specific row – or column if it is used on the right side of the comma – is removed. Then, the probability of suffering from coronary heart disease for observation i is predicted using this model and written into the previously for this purpose specified vector. For the prediction of an observations using an already fitted model we use the `predict()` command that need the model as an `glm` object and a set of the respective explanatory variables for these observations as input. Furthermore the `type` argument needs to be set. `response` means in this case that we want to get the output in terms of a fitted value, i.e. as a probability in case of logistic regression. To be more precise, the input of a `glm` object indicates that the function `predict.glm()` will be called that can handle regression model in particular. The `predict()` function can handle several different object types that are suitable for predictions. Hence, there are several parameters that may be set or passed additionally to this function. But just giving a specific object type as input suffices here to get the desired output, what makes it easy to

use in a lot of applications.

After the loop is done, the vector `prediction` contains the fitted values for each observation. Now, we check for each observation if the threshold given by `cutoff` is exceeded. If so, this yields a `TRUE` and otherwise a `FALSE` value. These values are in turn compared to the actual value of `chd`. Thus, we obtain the number of correct predictions by summing this up as well as the prediction accuracy by taking the proportion of all predictions. Returned is the prediction accuracy, in this case called the out-of-sample performance – since we are simulating to test the model on new data, as described in section 2.3.

model	out-of-sample performance
m0	.729
m2	.736
m3	.717

Table 7: The out-of-sample performance for three different models from leave-one-out cross-Validation (see related [Quantlet SAHeart.Q5.Leave-One-Out-Cross-Validation](#)).

The `loo.cv` function can now easily be used, simply inserting the models calculated in section 4.2. Table 7 shows the results. The model `m2` has the best out-of-sample performance. This is in line with the results of section 4.2 that were driven by regression analysis using the Akaike information criterion and significances of the coefficients as well as theoretical aspects. However, we still may improve our model.

4.4 Cook’s Distance

A further approach to improve the model is to find single observations that may influence the coefficients of the regression model in an inappropriate high manner. In section 3.2, we already saw some observations that are suspicious. E.g. the observation with the `id` 21 that is on the one hand very young, 17 years old, and has quite low values on all important risk factor, but on the other hand suffers from coronary heart disease, as can be seen from figure 6. This is very unlikely and the question arises if there happened a mistake during coding or if there are other reasons for that, e.g. some variables that are not represented in the dataset. Even if we say

that this might be a case that occurs with very low probability in the population and therefore might be represented in a random sample, the sample at hand is not that large, so that such occurrences have a larger impact on the estimates of the regression coefficients than we would assume for the population. In short, such inappropriate high impacts level out in large samples but can have too much weight due to small sample sizes. In this case this might lead to underestimating the influence of the single risk factors.

An easy way that helps out here is to exclude the respective observations. But one needs to be cautious. The other way around there is the risk to exclude too many observations that don't really fit the needs of the models and thus over-fitting the model to the given data. Or, as it is sometimes put, to fit the data to the model instead of fitting the model to the data.

A straightforward method that helps to detect such observations is the Cook's Distance. We wrote a function that calculates the Cook's Distance for each observation and plots the results to directly spot influential observations at first glance. The function `cookPlot` takes three arguments. As in section 4.3, `model` needs to be a `glm` or `lm` object. The optional argument `threshold` can be used to set the threshold above which the Cook's Distance values are denoted as influential or at least as suspiciously high. And the third parameter `PDFpath` is again used to directly save the plot to the local computer.

```

98 cookPlot <- function(model, threshold = NULL, PDFpath = NULL){
99   cd <- cooks.distance(m2)
100
101   if(!is.null(threshold)){
102     additional_params <- list(geom_text(data = data.frame(id = SAheart$id,
103                                                            cd = cd)[cd >= threshold,],
104                               aes(label = id), colour = 'firebrick',
105                               size = 3, vjust = -0.5),
106                               geom_hline(yintercept = threshold, color = 'firebrick'))
107   } else {
108     additional_params <- NULL
109   }
110
111   final.plot <- ggplot(data = data.frame(id = SAheart$id, cd = cd),
112                        mapping = aes(x = as.factor(id), y = cd)) +
113     geom_bar(stat = "identity", width = 0.1, color = "black") +
114     theme_classic() +
115     labs(x = "", y = "Cook's D") +
116     theme(axis.text.x = element_blank(), axis.ticks.x = element_blank()) +
117     scale_y_continuous(expand = c(0, 0), limits = c(0, max(cd)+0.002)) +
118     additional_params
119
120   plot(final.plot)
121
122   if(!is.null(PDFpath)){
123     ggsave(filename = "/CooksDplot.pdf", path = PDFpath, device = "pdf", plot = final.plot)
124   }
125 }

```

Code 16: Snippet of [Quantlet SAHeart_Q6.Cooks.Distance](#)

For the calculation of the Cook's Distance values we use the function `cooks.distance()` from the `stats`-package that is loaded by default. Like the `predict()` command from section 4.3 this function takes a `glm` or `lm` object as input and calls the respective function `cooks.distance.glm()` or `cooks.distance.lm()`, respectively.⁴ The output is a vector containing the Cook's Distance values for each observation that was given in the input model.

Furthermore, we check whether the `threshold` parameter has been specified using an `if` condition. A very nice feature of `ggplot()` is that we can append new layers and parameters by passing a list. Each element of the list will be handled as they were combined with `+`, what is the usual `ggplot()`-notation for adding layers and parameters. If this list is empty, i.e. `NULL`, it is just ignored. This makes it very easily possible, as it is done here, to add additional elements to the graphic on conditions. Since `ggplot()` is very flexible and comprehensive in parameters and graphic types there is a variety of possibilities to combine them with conditions. That makes it a remarkable tool for graphics. However, here we add two things to the graphic. First, a horizontal line marking the threshold with the `geom_hline` command and second the `id` of each observation that exceeds this threshold using the `geom_text` command. This can be seen in Figure 5. The data necessary for the labeling is here passed on the fly as a `data.frame` consisting of the `id` and the respective Cook's Distance values given that the value exceeds the threshold, using index notation. The `vjust` option, for *vertical adjustment*, adds some space between the bar and the label for a better appearance.

The commands to create the graphic have mostly been described above already. But notable is the `scale_y_continuous` command which is used to set the limits of the y-axis since we need some extra space at the top in case the labels are added. This is achieved by adding a little space to the maximum possible value, that is easily obtained with the function `max()`. And setting the `expand` parameter to `c(0,0)` removes the gap between the plot area and x-axis that is set by default. Also note how easy the list `additional_params` is added here.

⁴Note that in case of Gaussian regression models equation (10) can be expressed analytically and therefore calculated in a more efficient way. For the `glm` case an algorithm for approximation is used (for details see, Williams, 1987).




Figure 5: The Cook's Distance values for every observation. The red line marks the threshold above which the observations are denoted as influential. Those observations are labeled with their id (see related [Quantlet SAHeart.Q6-Cooks_Distance](#)).

The function doesn't return anything and just plots the graphic. To decide where the threshold should lie, we just ran `cookPlot` first without specifying the `threshold` parameter and obtained it as a visual estimate. Then, we run the function again with the threshold specified.

```

130 cookPlot(m2, threshold = 0.015, PDFpath = PDFpath)
131
132 cd <- cooks.distance(m2)
133
134 obslist <- SAheart$Id[cd >= 0.015]
135 varlist <- c("typea", "tobacco", "age", "ldl", "famhist")
136 multiJP(varlist, obslist, label = TRUE, PDFpath = PDFpath)

```

Code 17: Snippet of [Quantlet SAHeart.Q6-Cooks_Distance](#)

We can see from figure 5 that there are six suspicious observations, including the observations 21 and 261 that we already spotted above. Especially observation 17 seems to have a high influence. To not just naively remove these observations on first movement, we use the `multiJP` function from above to get some more insights on these observations.

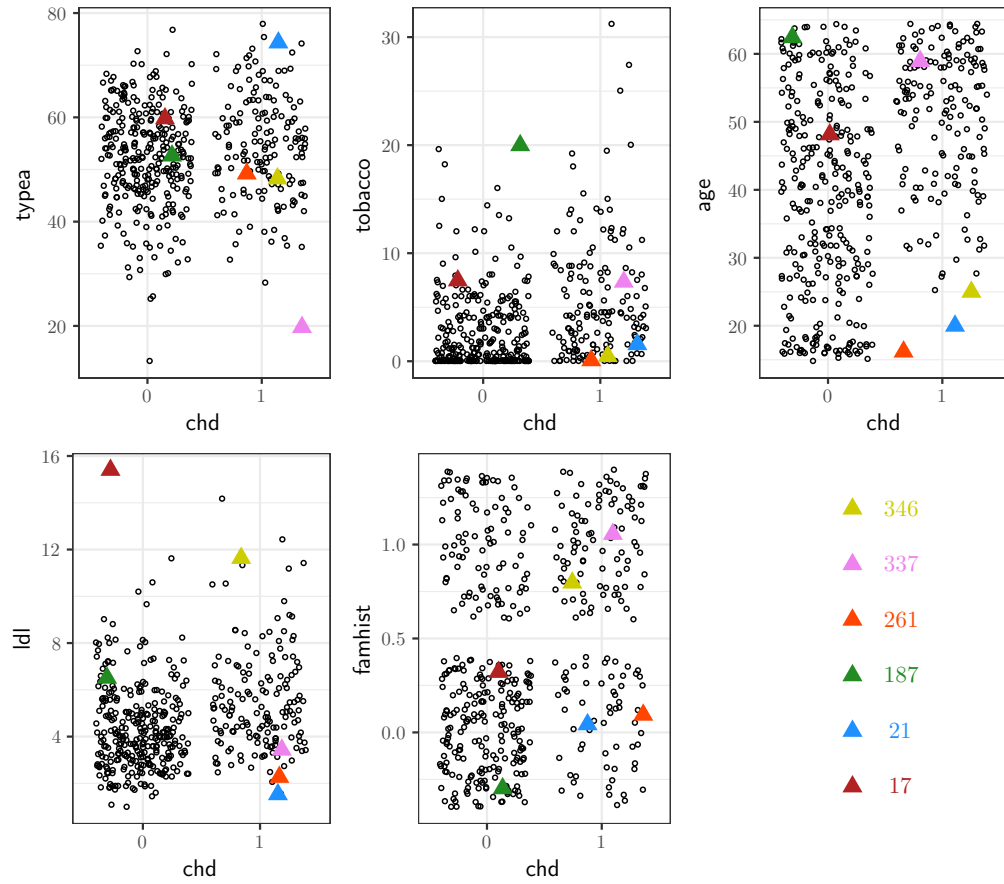


Figure 6: The distributions of the covariates over `chd`. Highlighted are the observations that have been detected as influential with Cook's Distance (see related [Quantlet SAHeart_Q6_Cooks_Distance](#)).

Figure 6 shows the behavior of these observation on the set of the explanatory variables we choose for our model. Observation 17 has high values – above the means – on `typea`, `tobacco` and `age` but no occurrence of coronary heart disease. This is not that conspicuous so far. But the value for `ldl` is very high compared to the other observations. As already mentioned above, we could ask if this is a mistake or just a rarely occurrence in the sample. However, for sake of not underestimating the single risk factors, what is obviously more fatal than overestimation – since in case of detection of diseases or risk factors for such, we rather take false positives than false negatives – we decided to exclude this specific observation. Also, for the same reasons, observations 21 and 261 are excluded. observation 261 shows very low values on all risk factors especially `age`. Observation 21 has also very low values on all these variables except for `typea`. Actually here, the value is pretty high. But

this variable is very controversial as already mentioned in section 3.1. Therefore, we still decided to drop it.

```

138 ## Drop all very influential obs for model calc
139 obslist <- SAheart$id[cd >= 0.015]
140 m2.1 <- glm(chd ~ tobacco + age + famhist + ldl + typea,
141             family = binomial(link = "logit"),
142             data = SAheart[!(SAheart$id %in% obslist),])
143 summary(m2.1)
144 loo_cv(m2.1)
145
146 # Drop only the influential points that are very suspicious in the graphics
147 # and whose excluding can be justified with theory
148 obslist <- c("17", "261", "21")
149 m2.2 <- glm(chd ~ tobacco + age + famhist + ldl + typea,
150             family = binomial(link = "logit"),
151             data = SAheart[!(SAheart$id %in% obslist),])
152 summary(m2.2)
153 loo_cv(m2.2)

```

Code 18: Snippet of [Quantlet SAHeart_Q6_Cooks_Distance](#)

Observation 187 is the most controversial case here. On the one hand, are the values for `tobacco` and `age` very high, but not abnormally high compared to other observations. On the other hand, the other variables do not really suggest the occurrence of coronary heart disease. So, getting old and smoking alone while not suffering from coronary heart disease is probably not that unusual. This is why we did not remove this observation. Also in the case of the observations 337 and 346 we didn't found them conspicuously enough to drop them.

In the next section we will use these insights to refit the models without these observations and evaluate the regression results as well as the out-of-sample performance from the `loo_cv` function.

4.5 Results

Throughout this paper we developed the model on theoretical reflections and graphical exploration resulting in a logistic regression model with a plausible set of explanatory variables, namely the model `m2`. Lastly, we spotted some observations that may be excluded from the dataset to refit the model without them. Thus, hopefully obtaining even better estimates. This is what we will do in this last section to select the final model that we would suggest from this dataset.

Table 8 shows the regression outputs for three models as well as the out-of-sample performance obtained from the `loo_cv` function. First, the model `m2` is presented here again for comparison reason. Model `m2.1` uses the same explanatory variables

	m2		m2.1		m2.2	
	coeff.	std. err.	coeff.	std. err.	coeff.	std. err.
Exclude observations	None		17, 21, 187, 261, 337, 346		17, 21, 261	
Coefficients						
constant	-6.446***	0.921	-7.241***	0.992	-6.853***	0.959
tobacco	0.080**	0.026	0.090***	0.027	0.088**	0.027
age	0.051***	0.010	0.057***	0.011	0.053***	0.011
famhist	0.908***	0.226	0.873***	0.232	0.923***	0.230
ldl	0.162**	0.055	0.196***	0.060	0.205***	0.059
typea	0.037**	0.012	0.042***	0.013	0.037**	0.012
AIC	487.690		464.500		472.150	
out-of-sample performance	0.736		0.745		0.747	

* $p < 0.05$, ** $p < 0.01$, *** $p < 0.001$

Table 8: Results of the Logistic Regression models fitted on different sets of observations. All models use the same set of explanatory variables as chosen for m2 in section 4.2 (see related [Quantlet SAHeart_Q4_Logistic_Regression](#), [Quantlet SAHeart_Q5_Leave-One-Out_Cross-Validation](#) and [Quantlet SAHeart_Q6_Cooks_Distance](#)) .

but excluding all observations that exceeded the threshold in Figure 5. And **m2.2** also uses the same explanatory variables but here all observations are excluded that could be excluded for theoretical reasons as discussed in section 4.4. The models **m2.1** and **m2.2** perform better on all indicators that we used for model selection here than model **m2**. The significance levels are better as well as especially the Akaike information criterion and the out-of-sample performance.

At first glance, one would select **m2.1** as the best model. But let us first reflect on what these measures mean and what they are based on. The significance level is here a measure of the certainty with which we can reject the null hypotheses that the true coefficient is zero. Of course, the higher this certainty the better. But the benefit gets smaller if we have already quite good p -values for all coefficients as it is the case here, actually at least $p < 0.01$. So, attaining even higher p -values from this start point is obviously not a sufficient logic here to select a model. Especially, and that is the most important point here, when this is attained by dropping observations that does not very well fit to the model. Because going on that way means from the perspective of equation (6) to successively remove the ϵ_i that are highest and thus, artificially shrinking the estimate of the variance $\hat{\sigma}_\epsilon^2$. Consequently, since the ϵ_i are the only stochastic components of the regression model, this affects the standard errors of the estimates and thus the p -values. And this is literally fitting the data to the model.

The same applies to the Akaike information criterion since it refers to the likelihood function of the model. And, for another reason, also to the out-of-sample performance since the `loo_cv` function only uses the observations that are passed through the input model. Therefore, the scope is still restricted to itself and that will automatically increase the prediction accuracy. However, to cut a long story short, the measures for model selection are of no value if theoretical aspects are ignored. Thus, unless model **m2.1** looks appealing, we choose **m2.2** as our final model since we can justify this theoretically as has been done in section 4.4.

To compare this to the Fisher's linear discriminant, where we achieved a prediction accuracy of 0.695, we managed to improve the prediction accuracy to 0.750, or by means of the out-of-sample performance, to 0.747. That is an increase of ca. 0.05. In

addition, the logistic regression model provides an overview of how strong the single variables effect the occurrence of coronary heart disease. This is a great advantage of the Fisher's linear discriminant method especially if the model is not just viewed from the perspective of classification but rather from the view of the risk factors itself.

5 Conclusion

As we already summarized the statistical results in the chapter before, the following will concentrate on our work with *R* and the real-world implications of the statistical results.

In the course of this paper we used different features of *R*. We loaded extensions and implemented code that deals with errors and shows easily understandable warning messages to the user. We analyzed data with the statistical tools of *R* and visualized the results in numerous complex graphics. Some of the code was written in form of functions, so that is usable for repeated cases by redefining the parameters. We also calculated the Fisher's linear discriminant from the scratch to show some fundamental functions of *R*, e.g. working with matrices.

Content-wise, our goal was to find the major risk factors of heart disease. Our descriptive analysis revealed, that there is a noteworthy correlation of **age**, **sbp**, **total-bacco**, **adiposity** and **ldl** levels with the occurrence of heart disease. The regression analysis gave further details, which factors how strongly influence the occurrence of heart disease. Using the leave-one-out cross-validation and Cook's distance, we determine model **m2.2** as our best model as it has the highest prediction accuracy. Its real-world implications are that behavioral factors like smoking and aggressive behavior play an important role, but there are also genetic factors like the family history (**famhist**) of heart disease. Additionally, cholesterol levels (**ldl**) and **age** are major predictors for the occurrence of heart disease. It is always dangerous to suggest concrete measures since the correlation doesn't necessarily implicates a causal relationship. It nevertheless seems to be a safe proposition that health organizations should intensify campaigns against smoking and should recommend

mandatory health checks for people at a certain age. These results are in line with the findings of Rossouw et al. (1983). We were not able to compare the detailed results as the data set in R only is a non-representative subset of the data Rossouw et al. (1983) used. In terms of conclusions, they additionally found that there should be health checks for people with obesity. As most participants exhibited at least one major risk factor, they even proposed to implement preventive measures for the whole community, which may be unrealistically costly for everything that goes further than information. Current research strongly suggests focusing on individuals with a high risk-score with the major factors tobacco, alcohol and overweight (obesity and adiposity) (cf. Putadechakum et al. (2014)).

References

- Cook, R. D. (1977). Detection of influential observations in linear regression. *Technometrics*, 19(1):15–18.
- Duda, R. O., Hart, P. E., and Stork, D. G. (2001). *Pattern classification*. Wiley & Sons, New York.
- Putadechakum, S., Leelahakul, V., Rodjinda, P., Phanachet, P., and Roongpisuthipong, C. (2014). Cardiovascular risk factors and 10-year risk for coronary heart disease in thai adults. www.ijsrp.org/research-paper-0514/ijsrp-p29119.pdf . Retrieved on March 24th, 2018.
- R-Foundation (2018). What is R? <https://www.r-project.org/about.html> . Retrieved on March 24th, 2018.
- Rossouw, J. E., Plessis, J. P. D., Benad, A. J. S., Jordaan, P. C. J., Kotz, J. P., Jooste, P. L., and Ferreira, J. J. (1983). Coronary risk factor screening in three rural communities – the coris baseline study. *South African Medical Journal*, 64.
- Spokoiny, V. and Dickhaus, T. (2015). *Basics of modern mathematical statistics*. Springer, Berlin/Heidelberg.
- Wasserman, L. (2013). *All of statistics: a concise course in statistical inference*. Springer Science & Business Media, New York.
- Wickham, H., Francois, R., Henry, L., Müller, K., and Wickham, M. H. (2017). Package ”dplyr”. <ftp://cran.de.r-project.org/pub/R/web/packages/dplyr/dplyr.pdf> . Retrieved on March 24th, 2018.
- Williams, D. A. (1987). Generalized linear model diagnostics using the deviance and single case deletions. *Applied Statistics*, 36:181–191.
- World Health Organization (2017). The top 10 causes of death - fact sheets. <http://www.who.int/mediacentre/factsheets/fs310/en/>. Retrieved on March 24th, 2018.

A Tables

	Estimate	Std. Error	z value	Pr(> z)	
(Intercept)	-6.1507	1.3083	-4.70	0.0000	***
sbp	0.0065	0.0057	1.14	0.2564	
tobacco	0.0794	0.0266	2.98	0.0028	**
ldl	0.1739	0.0597	2.92	0.0036	**
adiposity	0.0186	0.0293	0.63	0.5257	
famhist	0.9254	0.2279	4.06	0.0000	***
typea	0.0396	0.0123	3.21	0.0013	**
obesity	-0.0629	0.0442	-1.42	0.1551	
alcohol	0.0001	0.0045	0.03	0.9784	
age	0.0452	0.0121	3.73	0.0002	***
Significance codes	0 \Rightarrow ***	0.001 \Rightarrow **	0.01 \Rightarrow *		
AIC	492.14				

Table 9: The summary statistics of the logistic regression of **m0** (see related [Quantlet SAHeart_Q4_Logistic_Regression](#)).

	Estimate	Std. Error	z value	Pr(> z)	
(Intercept)	-6.3603	0.9356	-6.80	0.0000	***
tobacco	0.0804	0.0259	3.10	0.0019	**
age	0.0529	0.0113	4.67	0.0000	***
famhist	0.9097	0.2259	4.03	0.0001	***
ldl	0.1724	0.0592	2.91	0.0036	**
typea	0.0371	0.0122	3.05	0.0023	**
adiposity	-0.0094	0.0192	-0.49	0.6229	
Significance codes	0 \Rightarrow ***	0.001 \Rightarrow **	0.01 \Rightarrow *		
AIC	489.44				

Table 10: The summary statistics of the logistic regression of **m1.1** (see related [Quantlet SAHeart_Q4_Logistic_Regression](#)).

	Estimate	Std. Error	z value	Pr(> z)	
(Intercept)	-5.7027	1.0764	-5.30	0.0000	***
tobacco	0.0800	0.0260	3.08	0.0021	**
age	0.0521	0.0102	5.09	0.0000	***
famhist	0.9161	0.2264	4.05	0.0001	***
ldl	0.1837	0.0582	3.16	0.0016	**
typea	0.0383	0.0122	3.13	0.0017	**
obesity	-0.0376	0.0291	-1.29	0.1964	
Significance codes	0 ⇒ ***	0.001 ⇒ **	0.01 ⇒ *		
AIC	487.98				

Table 11: The summary statistics of the logistic regression of `m1.2` (see related [Quantlet SAHeart_Q4_Logistic_Regression](#)).

	Estimate	Std. Error	z value	Pr(> z)	
(Intercept)	-4.2043	0.4983	-8.44	0.0000	***
tobacco	0.0807	0.0255	3.16	0.0016	**
age	0.0440	0.0097	4.52	0.0000	***
famhist	0.9241	0.2232	4.14	0.0000	***
ldl	0.1676	0.0542	3.09	0.0020	**
Significance codes	0 ⇒ ***	0.001 ⇒ **	0.01 ⇒ *		
AIC	495.44				

Table 12: The summary statistics of the logistic regression of `m3` (see related [Quantlet SAHeart_Q4_Logistic_Regression](#)).

B Code

B.1 Descriptive Statistics

B.1.1 Quantlet 1 - Data Preparation

```

1 ## Check if required packages are installed, if not install and load
2 loadPKG <- function(pkg){
3   if (!(pkg %in% installed.packages())){
4     message(paste0("The required package '",
5                     pkg, "' is currently not installed. ",
6                     "It will now be installed to the default library path. "))
7     install.packages(pkg)
8   }
9   library(pkg, character.only = TRUE)
10 }
11
12 loadPKG("dplyr")
13 loadPKG("ggplot2")
14 loadPKG("gridExtra")
15 loadPKG("ElemStatLearn")
16 loadPKG("corrplot")
17 loadPKG("xtable")

```

```

18 |
19 | data(SAheart)
20 |
21 |
22 | ## Specify Path to save graphic outputs
23 | PDFpath <- "~/Desktop"
24 |
25 |
26 | ## check if PDFpath exists, if not choose current working directory
27 | if (!dir.exists(PDFpath)){
28 |   PDFpath <- getwd()
29 |   message(paste0("The \'PDFpath\' that was specified does not exist. ",
30 |                 "Instead, the graphic outputs will be saved to the current ",
31 |                 "working directory. The current working directory is: \'",
32 |                 getwd(), "\'"))
33 | }
34 |
35 |
36 | #####
37 |
38 |
39 | ## Variables of dataset
40 | # sbp          systolic blood pressure
41 | #
42 | # tobacco      cumulative tobacco (kg)
43 | #
44 | # ldl          low density lipoprotein cholesterol
45 | #
46 | # adiposity    a numeric vector
47 | #
48 | # famhist      family history of heart disease, a factor with levels Absent Present
49 | #
50 | # typea        type-A behavior
51 | #
52 | # obesity      a numeric vector
53 | #
54 | # alcohol      current alcohol consumption
55 | #
56 | # age          age at onset
57 | #
58 | # chd          response, coronary heart disease
59 |
60 |
61 | #####
62 | ##              ##
63 | ## Description of dataset ##
64 | ##              ##
65 | #####
66 |
67 |
68 | ## Some overview stuff
69 |
70 | glimpse(SAheart)
71 | str(SAheart)
72 |
73 | head(SAheart, 5)
74 |
75 | summary(SAheart)
76 |
77 |
78 | ## Create id variable to easier handle single observations (using 'dplyr')
79 | SAheart <- SAheart %>%
80 |   mutate(id = row_number())
81 |
82 |
83 | ## famhist
84 | head(SAheart$famhist, 10)
85 | head(as.numeric(SAheart$famhist), 10)
86 | SAheart$famhist <- as.numeric(SAheart$famhist)
87 | SAheart$famhist[SAheart$famhist == 1] <- 0
88 | SAheart$famhist[SAheart$famhist == 2] <- 1

```

Code 19:  Quantlet SAHeart_Q1_Data_Preparation

B.1.2 Quantlet 2 - Descriptive Statistics

```
1
2 ## Check if required packages are installed, if not install and load
3 loadPKG <- function(pkg){
4   if (!(pkg %in% installed.packages())){
5     message(paste0("The required package \'",
6                     pkg, "\' is currently not installed. ",
7                     "It will now be installed to the default library path. "))
8     install.packages(pkg)
9   }
10  library(pkg, character.only = TRUE)
11 }
12
13 loadPKG("dplyr")
14 loadPKG("ggplot2")
15 loadPKG("gridExtra")
16 loadPKG("ElemStatLearn")
17 loadPKG("corrplot")
18 loadPKG("xtable")
19
20 data(SAheart)
21
22
23 ## Specify Path to save graphic outputs
24 PDFpath <- "~/Desktop"
25
26
27 ## check if PDFpath exists, if not choose current working directory
28 if(!dir.exists(PDFpath)){
29   PDFpath <- getwd()
30   message(paste0("The \\'PDFpath\\' that was specified does not exist. ",
31                   "Instead, the graphic outputs will be saved to the current ",
32                   "working directory. The current working directory is: \'",
33                   getwd(), "\'"))
34 }
35
36
37
38 ## Create id variable to easier handle single observations (using 'dplyr')
39 SAheart <- SAheart %>%
40   mutate(id = row_number())
41
42
43 ## famhist
44 head(SAheart$famhist, 10)
45 head(as.numeric(SAheart$famhist), 10)
46 SAheart$famhist <- as.numeric(SAheart$famhist)
47 SAheart$famhist[SAheart$famhist == 1] <- 0
48 SAheart$famhist[SAheart$famhist == 2] <- 1
49
50
51
52 #####
53 ##           Exploratory Analysis           ##
54 ##                                           ##
55 ##                                           ##
56 #####
57
58
59 ## Correlation matrix (using package corrplot)
60 pdf(file = paste0(PDFpath, "/Corrplot.pdf"), width = 7, height = 8)
61 corr_matrix <- cor(select(SAheart, chd, everything()), -id)
62 colset <- colorRampPalette(c("#BB4444", "#EE9988", "#FFFFFF", "#77AADD", "#4477AA"))
63 corrplot.mixed(corr_matrix,
64                upper = "ellipse",
65                lower = "number",
66                number.cex = .7,
67                upper.col = colset(10),
68                lower.col = "black",
69                tl.col = "black",
70                tl.pos = "lt")
71 dev.off()
72
73
74 ## Function to show Boxplots of covariates over chd
75 multiBP <- function(varlist, PDFpath = NULL){
76   plot.list <- lapply(varlist, function(ivar){
77     ggplot(data = SAheart, mapping = aes(x = as.factor(chd), y = get(ivar))) +
78       geom_boxplot(stat = "boxplot", position = "dodge") +
79       theme_bw() +
80       xlab("chd") +
81       ylab(ivar)
82   })
83
84   if(length(plot.list) <= 3){
85     ncol <- length(plot.list)
86   } else {
87     ncol <- 3
88   }
89
90   final.plot <- grid.arrange(grobs = plot.list, ncol=ncol)
91
92   if (!is.null(PDFpath)){
```

```

93 |     ggsave(filename = "/BPplot.pdf", path = PDFpath, device = "pdf", plot = final.plot)
94 | }
95 | }
96 |
97 |
98 | varlist <- c("sbp", "tobacco", "age", "ldl", "adiposity")
99 | multiBP(varlist, PDFpath)
100 |
101 |
102 | ## Suspicious outlier
103 | SAheart$id[SAheart$age <= 20 & SAheart$chd == 1]
104 |
105 | subset(SAheart, select = varlist, id == 261)
106 | subset(SAheart, select = varlist, id == 21)
107 | summary(subset(SAheart, select = varlist, chd == 1))
108 |
109 | ## Function to visualize behaviour of single observations
110 | multiJP <- function(varlist, obslist = 0, PDFpath = NULL, label = FALSE){
111 |
112 |
113 |   if(label){
114 |     varlist <- c(varlist, "legend")
115 |   }
116 |
117 |   color.list <- c("firebrick", "dodgerblue", "forestgreen",
118 |                  "orangered", "violet", "yellow3")
119 |
120 |   plot.list <- lapply(varlist, function(ivar){
121 |     if(ivar == "legend"){
122 |       hackdata <- data.frame(id = obslist,
123 |                             position = seq(1, length(obslist)),
124 |                             xvalue = 0.11)
125 |       legend <- ggplot(data = hackdata, mapping = aes(x = xvalue, y = position)) +
126 |         geom_point(colour = color.list[1:length(obslist)], shape = 17, size = 3) +
127 |         theme_classic() +
128 |         theme(axis.text = element_blank(), axis.ticks = element_blank(),
129 |               axis.line = element_blank()) +
130 |         labs(x = "", y = "") +
131 |         scale_y_continuous(expand = c(0, 0), limits = c(0, max(hackdata$position)+1)) +
132 |         scale_x_continuous(limits = c(0.1, 1)) +
133 |         geom_text(data = hackdata, aes(label = id),
134 |                   colour = color.list[1:length(obslist)], size = 3, hjust = -1)
135 |
136 |     } else {
137 |
138 |       ggplot(data = SAheart[!(SAheart$id %in% obslist),],
139 |             mapping = aes(x = as.factor(chd), y = get(ivar))) +
140 |         geom_jitter(shape = 1) +
141 |         theme_bw() +
142 |         xlab("chd") +
143 |         ylab(ivar) +
144 |         geom_jitter(data = SAheart[SAheart$id %in% obslist,],
145 |                     colour = color.list[1:length(obslist)], shape = 17, size = 3) # +
146 |     }
147 |   })
148 |
149 |   if(length(plot.list) <= 3){
150 |     ncol <- length(plot.list)
151 |   } else {
152 |     ncol <- 3
153 |   }
154 |
155 |   final.plot <- grid.arrange(grobs = plot.list, ncol=ncol)
156 |
157 |   if(!is.null(PDFpath)){
158 |     ggsave(filename = "/Jitterplot-outliers.pdf", path = PDFpath, device = "pdf", plot = final.plot)}
159 |
160 |   varlist <- c("sbp", "tobacco", "age", "ldl", "adiposity")
161 |   multiJP(varlist, obslist = c(261))
162 |
163 |
164 |   ## Also check for the other very young person:
165 |   multiJP(varlist, obslist = c(261, 21), label = TRUE, PDFpath = PDFpath)
166 |
167 |
168 |   ## Check in general for obs that have chd but have very low values on all important variables
169 |   obslist <- (SAheart %>%
170 |     filter(chd == 1,
171 |            sbp <= quantile(sbp, .25),
172 |            tobacco <= quantile(tobacco, .25),
173 |            age <= quantile(age, .25),
174 |            ldl <= quantile(ldl, .25),
175 |            adiposity <= quantile(adiposity, .25),
176 |            famhist == 0) %>%
177 |     select(id))$id

```

Code 20: [Quantlet SAHeart-Q2.DescriptiveStatistics](#)

B.2 Inferential Statistics

B.2.1 Quantlet 3 - Fisher's Linear Discriminant

```
1
2 ## Check if required packages are installed, if not install and load
3 loadPKG <- function(pkg){
4   if (!(pkg %in% installed.packages())){
5     message(paste0("The required package '",
6                     pkg, "' is currently not installed. ",
7                     "It will now be installed to the default library path."))
8     install.packages(pkg)
9   }
10  library(pkg, character.only = TRUE)
11 }
12
13 loadPKG("dplyr")
14 loadPKG("ggplot2")
15 loadPKG("gridExtra")
16 loadPKG("ElemStatLearn")
17 loadPKG("corrplot")
18 loadPKG("xtable")
19
20 data(SAheart)
21
22
23 ## Specify Path to save graphic outputs
24 PDFpath <- "~/Desktop"
25
26
27 ## check if PDFpath exists, if not choose current working directory
28 if (!dir.exists(PDFpath)){
29   PDFpath <- getwd()
30   message(paste0("The 'PDFpath' that was specified does not exist. ",
31                   "Instead, the graphic outputs will be saved to the current ",
32                   "working directory. The current working directory is: '",
33                   getwd(), "'"))
34 }
35
36 ## Create id variable to easier handle single observations (using 'dplyr')
37 SAheart <- SAheart %>%
38   mutate(id = row.number())
39 SAheart$famhist <- as.numeric(SAheart$famhist)
40 SAheart$famhist[SAheart$famhist == 1] <- 0
41 SAheart$famhist[SAheart$famhist == 2] <- 1
42
43
44 #####
45 ## ##
46 ## Fisher Linear Discriminant ##
47 ## ##
48 #####
49
50
51 ## Function for Fisher Linear Discriminant
52 fisher <- function(X, c, PDFpath = NULL){
53
54   X <- as.matrix(X)
55
56   X0 <- X[c == 0,]
57   X1 <- X[c == 1,]
58
59   n0 <- nrow(X0)
60   n1 <- nrow(X1)
61
62   m0 <- as.vector(colMeans(X0))
63   m1 <- as.vector(colMeans(X1))
64
65   M0 <- as.vector(rep(1, n0)) %*% t(m0)
66   M1 <- as.vector(rep(1, n1)) %*% t(m1)
67
68   X0_c <- X0 - M0
69   X1_c <- X1 - M1
70
71   S0 <- t(X0_c) %*% X0_c
72   S1 <- t(X1_c) %*% X1_c
73
74   S_w <- S0 + S1
75
76   S_w_inv <- solve(S_w)
77
78   w <- S_w_inv %*% (m0 - m1)
79
80   m <- as.vector(colMeans(X))
81   M <- as.vector(rep(1, nrow(X))) %*% t(m)
82   X_c <- X - M
83
84   z <- t(w) %*% t(X_c) < 0
85
86   accuracy <- sum(t(z) == c) / length(z)
87
88   Y0 <- X0 %*% w
```



```

89 | Y1 <- X1 %*% w
90 |
91 | Y <- rbind(cbind(Y0, rep(0,nrow(Y0))),
92 |           cbind(Y1, rep(1,nrow(Y1))))
93 | Y <- as.data.frame(Y)
94 | colnames(Y) <- c("Y", "c")
95 |
96 | final.plot <- ggplot(Y, aes(Y, fill = as.factor(c))) + geom_density(alpha = 0.2) +
97 |   theme_classic() +
98 |   labs(x = "", y = "") +
99 |   theme(axis.text.x = element_blank(), axis.ticks.x = element_blank(),
100 |         axis.text.y = element_blank(), axis.ticks.y = element_blank(),
101 |         axis.line.y = element_blank()) +
102 |     scale_fill_manual(values = c("firebrick", "dodgerblue"),
103 |                       name = "chd", labels = c("no", "yes"))
104 |
105 | plot(final.plot)
106 |
107 | if(!is.null(PDFpath)){
108 |   ggsave(filename = "/Densitiesplot.pdf", path = PDFpath, device = "pdf", plot = final.plot)
109 | }
110 |
111 | return(accuracy)
112 | }
113 |
114 |
115 | X <- select(SAheart, -id, -chd)
116 | c <- SAheart$chd
117 |
118 | fisher(X=X, c=c, PDFpath)

```

Code 21: [Quantlet SAHeart_Q3_Fisher_Linear_Discriminant](#)

B.2.2 Quantlet 4 - Logistic Regression

```

1 |
2 | ## Check if required packages are installed, if not install and load
3 | loadPKG <- function(pkg){
4 |   if (!(pkg %in% installed.packages())){
5 |     message(paste0("The required package '",
6 |                   pkg, "' is currently not installed. ",
7 |                   "It will now be installed to the default library path."))
8 |     install.packages(pkg)
9 |   }
10 |   library(pkg, character.only = TRUE)
11 | }
12 |
13 | loadPKG("dplyr")
14 | loadPKG("ggplot2")
15 | loadPKG("gridExtra")
16 | loadPKG("ElemStatLearn")
17 | loadPKG("corrplot")
18 | loadPKG("xtable")
19 |
20 | data(SAheart)
21 |
22 |
23 | ## Specify Path to save graphic outputs
24 | PDFpath <- "~/Desktop"
25 |
26 |
27 | ## check if PDFpath exists, if not choose current working directory
28 | if(!dir.exists(PDFpath)){
29 |   PDFpath <- getwd()
30 |   message(paste0("The 'PDFpath' that was specified does not exist. ",
31 |                 "Instead, the graphic outputs will be saved to the current ",
32 |                 "working directory. The current working directory is: '",
33 |                 getwd(), "'"))
34 | }
35 |
36 |
37 |
38 | ## Create id variable to easier handle single observations (using 'dplyr')
39 | SAheart <- SAheart %>%
40 |   mutate(id = row_number())
41 | SAheart$famhist <- as.numeric(SAheart$famhist)
42 | SAheart$famhist[SAheart$famhist == 1] <- 0
43 | SAheart$famhist[SAheart$famhist == 2] <- 1
44 |
45 | ## Function to show Boxplots of covariates over chd
46 | multiBP <- function(varlist, PDFpath = NULL){
47 |   plot.list <- lapply(varlist, function(ivar){
48 |     ggplot(data = SAheart, mapping = aes(x = as.factor(chd), y = get(ivar))) +
49 |       geom_boxplot(stat = "boxplot", position = "dodge") +
50 |       theme_bw() +
51 |       xlab("chd") +
52 |       ylab(ivar)
53 |   })

```

```

54 |
55 |   if (length(plot.list) <= 3){
56 |     ncol <- length(plot.list)
57 |   } else{
58 |     ncol <- 3
59 |   }
60 |
61 |   final.plot <- grid.arrange(grobs = plot.list, ncol=ncol)
62 | }
63 |
64 | ## Function to visualize behaviour of single observations
65 | multiJP <- function(varlist, obslist = 0, PDFpath = NULL, label = FALSE){
66 |
67 |   if (label){
68 |     labellist <- obslist
69 |     varlist <- c(varlist, "legend")
70 |   } else{
71 |     labellist <- ""
72 |   }
73 |
74 |   color.list <- c("firebrick", "dodgerblue", "forestgreen",
75 |                  "orangered", "violet", "yellow3")
76 |
77 |   plot.list <- lapply(varlist, function(ivar){
78 |     if (ivar == "legend"){
79 |       hackdata <- data.frame(id = obslist,
80 |                              position = seq(1, length(obslist)),
81 |                              xvalue = 0.11)
82 |       legend <- ggplot(data = hackdata, mapping = aes(x = xvalue, y = position)) +
83 |         geom_point(colour = color.list[1:length(obslist)], shape = 17, size = 3) +
84 |         theme_classic() +
85 |         theme(axis.text = element_blank(), axis.ticks = element_blank(),
86 |               axis.line = element_blank()) +
87 |         labs(x = "", y = "") +
88 |         scale_y_continuous(expand = c(0, 0), limits = c(0, max(hackdata$position)+1)) +
89 |         scale_x_continuous(limits = c(0.1, 1)) +
90 |         geom_text(data = hackdata, aes(label = id),
91 |                  colour = color.list[1:length(obslist)], size = 3, hjust = -1)
92 |     } else{
93 |
94 |       ggplot(data = SAheart[!(SAheart$id %in% obslist),],
95 |             mapping = aes(x = as.factor(chd), y = get(ivar))) +
96 |         geom_jitter(shape = 1) +
97 |         theme_bw() +
98 |         xlab("chd") +
99 |         ylab(ivar) +
100 |         geom_jitter(data = SAheart[SAheart$id %in% obslist,],
101 |                    colour = color.list[1:length(obslist)], shape = 17, size = 3) # +
102 |     }
103 |   })
104 |
105 |   if (length(plot.list) <= 3){
106 |     ncol <- length(plot.list)
107 |   } else{
108 |     ncol <- 3
109 |   }
110 |   final.plot <- grid.arrange(grobs = plot.list, ncol=ncol)
111 | }
112 |
113 |
114 | #####
115 | ##                               ##
116 | ##      Logistic Regression      ##
117 | ##                               ##
118 | #####
119 |
120 |
121 | ## M0: (naive) Full model
122 | m0 <- glm(chd ~ ., family = binomial(link = "logit"),
123 |          data = select(SAheart, -id))
124 | summary(m0)
125 |
126 |
127 | ## M1: Without sbp, alcohol; and exclude either adiposity or obesity
128 | m1.1 <- glm(chd ~ tobacco + age + famhist + ldl + typea + adiposity,
129 |            family = binomial(link = "logit"), data = SAheart)
130 | summary(m1.1)
131 |
132 | m1.2 <- glm(chd ~ tobacco + age + famhist + ldl + typea + obesity,
133 |            family = binomial(link = "logit"), data = SAheart)
134 | summary(m1.2)
135 |
136 |
137 |
138 | ## M2: Exclude sbp, alcohol, adiposity, obesity
139 | m2 <- glm(chd ~ tobacco + age + famhist + ldl + typea,
140 |          family = binomial(link = "logit"), data = SAheart)
141 | summary(m2)
142 |
143 |
144 | ## M3: For sake of simplicity of models we try to exclude
145 | ## typea, since effect seems to be still very low
146 | m3 <- glm(chd ~ tobacco + age + famhist + ldl,
147 |          family = binomial(link = "logit"), data = SAheart)
148 | summary(m3)

```

```

149
150
151
152 ## typea: mediator effect via age
153
154 mean(SAheart$chd[SAheart$typea >= median(SAheart$typea)])
155 mean(SAheart$chd[SAheart$typea < median(SAheart$typea)])
156
157 multiBP("typea")
158
159 obslist <- (SAheart %>%
160   filter(chd == 1 & typea < 30) %>%
161   select(id))$id
162
163 varlist <- c("sbp", "tobacco", "age", "ldl", "typea", "famhist")
164 multiJP(varlist, obslist, label = TRUE)
165
166
167 summary(glm(chd ~ typea, family = binomial(link = "logit"), data = SAheart))
168 summary(glm(chd ~ typea + age, family = binomial(link = "logit"), data = SAheart))

```

Code 22: [Quantlet SAHeart_Q4_Logistic_Regression](#)

B.2.3 Quantlet 5 - Leave-One-Out Cross-Validation

```

1
2 ## Check if required packages are installed, if not install and load
3 loadPKG <- function(pkg){
4   if (!(pkg %in% installed.packages())){
5     message(paste0("The required package '",
6       pkg, "' is currently not installed. ",
7       "It will now be installed to the default library path."))
8     install.packages(pkg)
9   }
10  library(pkg, character.only = TRUE)
11 }
12
13 loadPKG("dplyr")
14 loadPKG("ggplot2")
15 loadPKG("gridExtra")
16 loadPKG("ElemStatLearn")
17 loadPKG("corrplot")
18 loadPKG("xtable")
19
20 data(SAheart)
21
22
23 ## Specify Path to save graphic outputs
24 PDFpath <- "~/Desktop"
25
26
27 ## check if PDFpath exists, if not choose current working directory
28 if (!dir.exists(PDFpath)){
29   PDFpath <- getwd()
30   message(paste0("The 'PDFpath' that was specified does not exist. ",
31     "Instead, the graphic outputs will be saved to the current ",
32     "working directory. The current working directory is: '",
33     getwd(), "'"))
34 }
35
36
37 ## Create id variable to easier handle single observations (using 'dplyr')
38 SAheart <- SAheart %>%
39   mutate(id = row_number())
40 SAheart$famhist <- as.numeric(SAheart$famhist)
41 SAheart$famhist[SAheart$famhist == 1] <- 0
42 SAheart$famhist[SAheart$famhist == 2] <- 1
43
44
45 ## M0: (naive) Full model
46 m0 <- glm(chd ~ ., family = binomial(link = "logit"),
47   data = select(SAheart, -tobacco, -id))
48
49 ## M1: Without sbp, alcohol; and exclude either adiposity or obesity
50 m1.1 <- glm(chd ~ tobacco + age + famhist + ldl + typea + adiposity,
51   family = binomial(link = "logit"), data = SAheart)
52
53 m1.2 <- glm(chd ~ tobacco + age + famhist + ldl + typea + obesity,
54   family = binomial(link = "logit"), data = SAheart)
55
56 ## M2: Exclude sbp, alcohol, adiposity, obesity
57 m2 <- glm(chd ~ tobacco + age + famhist + ldl + typea,
58   family = binomial(link = "logit"), data = SAheart)
59
60 ## M3: For sake of simplicity of models we try to exclude typea, since effect seems to be still very low
61 m3 <- glm(chd ~ tobacco + age + famhist + ldl,
62   family = binomial(link = "logit"), data = SAheart)
63

```

```

64 #####
65 ##                                ##
66 ##          loo-cv function          ##
67 ##                                ##
68 ##                                ##
69 #####
70
71
72 loo_cv <- function(model, cutoff = 0.5){
73
74   cur.data <- model$data
75   cur.formula <- model$formula
76   N <- nrow(cur.data)
77
78   prediction <- vector(mode = "numeric", length = N)
79
80   for(i in 1:N){
81     cur.model <- glm(cur.formula, family = binomial(link = "logit"),
82                      data = cur.data[-i,])
83     prediction[i] <- predict(cur.model, newdata = cur.data[i,], type = "response")
84   }
85
86   correct <- cur.data$chd == (prediction > cutoff)
87
88   accuracy <- sum(correct)/N
89
90   return(accuracy)
91 }
92
93
94 loo_cv(m0)
95 loo_cv(m2)
96 loo_cv(m3)

```

Code 23: Quantlet SAHeart_Q5_Leave-One-Out_Cross-Validation

B.2.4 Quantlet 6 - Cook's Distance

```

1
2 ## Check if required packages are installed, if not install and load
3 loadPKG <- function(pkg){
4   if (!(pkg %in% installed.packages())){
5     message(paste0("The required package '",
6                     pkg, "' is currently not installed. ",
7                     "It will now be installed to the default library path. "))
8     install.packages(pkg)
9   }
10  library(pkg, character.only = TRUE)
11 }
12
13 loadPKG("dplyr")
14 loadPKG("ggplot2")
15 loadPKG("gridExtra")
16 loadPKG("ElemStatLearn")
17 loadPKG("corrplot")
18 loadPKG("xtable")
19
20 data(SAheart)
21
22
23 ## Specify Path to save graphic outputs
24 PDFpath <- "~/Desktop"
25
26
27 ## check if PDFpath exists, if not choose current working directory
28 if (!dir.exists(PDFpath)){
29   PDFpath <- getwd()
30   message(paste0("The 'PDFpath' that was specified does not exist. ",
31                   "Instead, the graphic outputs will be saved to the current ",
32                   "working directory. The current working directory is: '",
33                   getwd(), "'"))
34 }
35
36
37
38 ## Create id variable to easier handle single observations (using 'dplyr')
39 SAheart <- SAheart %>%
40   mutate(id = row.number())
41 SAheart$famhist <- as.numeric(SAheart$famhist)
42 SAheart$famhist[SAheart$famhist == 1] <- 0
43 SAheart$famhist[SAheart$famhist == 2] <- 1
44
45
46
47 ## M0: (naive) Full model
48 m0 <- glm(chd ~ ., family = binomial(link = "logit"),
49           data = select(SAheart, -tobacco, -id))

```

```

50 |
51 |
52 | ## M1: Without sbp, alcohol; and exclude either adiposity or obesity
53 | m1.1 <- glm(chd ~ tobacco + age + famhist + ldl + typea + adiposity,
54 |           family = binomial(link = "logit"), data = SAheart)
55 |
56 |
57 | m1.2 <- glm(chd ~ tobacco + age + famhist + ldl + typea + obesity,
58 |           family = binomial(link = "logit"), data = SAheart)
59 |
60 |
61 | ## M2: Exclude sbp, alcohol, adiposity, obesity
62 | m2 <- glm(chd ~ tobacco + age + famhist + ldl + typea,
63 |           family = binomial(link = "logit"), data = SAheart)
64 |
65 |
66 | ## M3: For sake of simplicity of models we try to exclude typea, since effect seems to be still very low
67 | m3 <- glm(chd ~ tobacco + age + famhist + ldl,
68 |           family = binomial(link = "logit"), data = SAheart)
69 |
70 | loo_cv <- function(model, cutoff = 0.5){
71 |
72 |   cur.data <- model$data
73 |   cur.formula <- model$formula
74 |   N <- nrow(cur.data)
75 |
76 |   prediction <- vector(mode = "numeric", length = N)
77 |
78 |   for(i in 1:N){
79 |     cur.model <- glm(cur.formula, family = binomial(link = "logit"),
80 |                     data = cur.data[-i,])
81 |     prediction[i] <- predict(cur.model, newdata = cur.data[i,], type = "response")
82 |   }
83 |
84 |   correct <- cur.data$chd == (prediction > cutoff)
85 |
86 |   accuracy <- sum(correct)/N
87 |
88 |   return(accuracy)
89 | }
90 |
91 | #####
92 | ##                                ##
93 | ##      Cook's Distance          ##
94 | ##                                ##
95 | #####
96 |
97 |
98 | cookPlot <- function(model, threshold = NULL, PDFpath = NULL){
99 |
100 |   cd <- cooks.distance(m2)
101 |
102 |   if(!is.null(threshold)){
103 |     additional_params <- list(geom_text(data = data.frame(id = SAheart$id,
104 |                                                           cd = cd)[cd >= threshold,],
105 |                               aes(label = id), colour = 'firebrick',
106 |                               size = 3, vjust = -0.5),
107 |                               geom_hline(yintercept = threshold, color = 'firebrick'))
108 |   } else {
109 |     additional_params <- NULL
110 |   }
111 |
112 |   final.plot <- ggplot(data = data.frame(id = SAheart$id, cd = cd),
113 |                        mapping = aes(x = as.factor(id), y = cd)) +
114 |     geom_bar(stat = "identity", width = 0.1, color = "black") +
115 |     theme_classic() +
116 |     labs(x = "", y = "Cook's D") +
117 |     theme(axis.text.x = element_blank(), axis.ticks.x = element_blank()) +
118 |     scale_y_continuous(expand = c(0, 0), limits = c(0, max(cd)+0.002)) +
119 |     additional_params
120 |
121 |   plot(final.plot)
122 |
123 |   if(!is.null(PDFpath)){
124 |     ggsave(filename = "/CooksDplot.pdf", path = PDFpath, device = "pdf", plot = final.plot)
125 |   }
126 |
127 | }
128 |
129 |
130 | cookPlot(m2, threshold = 0.015, PDFpath = PDFpath)
131 |
132 | cd <- cooks.distance(m2)
133 |
134 | obslist <- SAheart$id[cd >= 0.015]
135 | varlist <- c("typea", "tobacco", "age", "ldl", "famhist")
136 | multiJP(varlist, obslist, label = TRUE, PDFpath = PDFpath)
137 |
138 | ## Drop all very influential obs for model calc
139 | obslist <- SAheart$id[cd >= 0.015]
140 | m2.1 <- glm(chd ~ tobacco + age + famhist + ldl + typea,
141 |            family = binomial(link = "logit"),
142 |            data = SAheart[!(SAheart$id %in% obslist),])
143 | summary(m2.1)
144 | loo_cv(m2.1)

```

```

145
146 # Drop only the influential points that are very suspicious in the graphics
147 # and whose excluding can be justified with theory
148 obslist <- c("17", "261", "21")
149 m2.2 <- glm(chd ~ tobacco + age + famhist + ldl + typea,
150             family = binomial(link = "logit"),
151             data = SAheart[!(SAheart$id %in% obslist),])
152 summary(m2.2)
153 loo_cv(m2.2)

```

Code 24: uantlet [SAHeart_Q6_Cooks_Distance](#)

C Declaration Of Authorship

We hereby confirm that we have authored this Seminar paper independently and without use of others than the indicated sources. All passages which are literally or in general matter taken out of publications or other sources are marked as such.

Berlin, March 30th 2018

Lukas Bargel, Adrian Rolf and Felix Vala

Three handwritten signatures in blue ink are displayed horizontally. From left to right: the first signature is 'Lukas Bargel', the second is 'Adrian Rolf' (written as 'Rolf'), and the third is 'Felix Vala' (written as 'F. Vala').